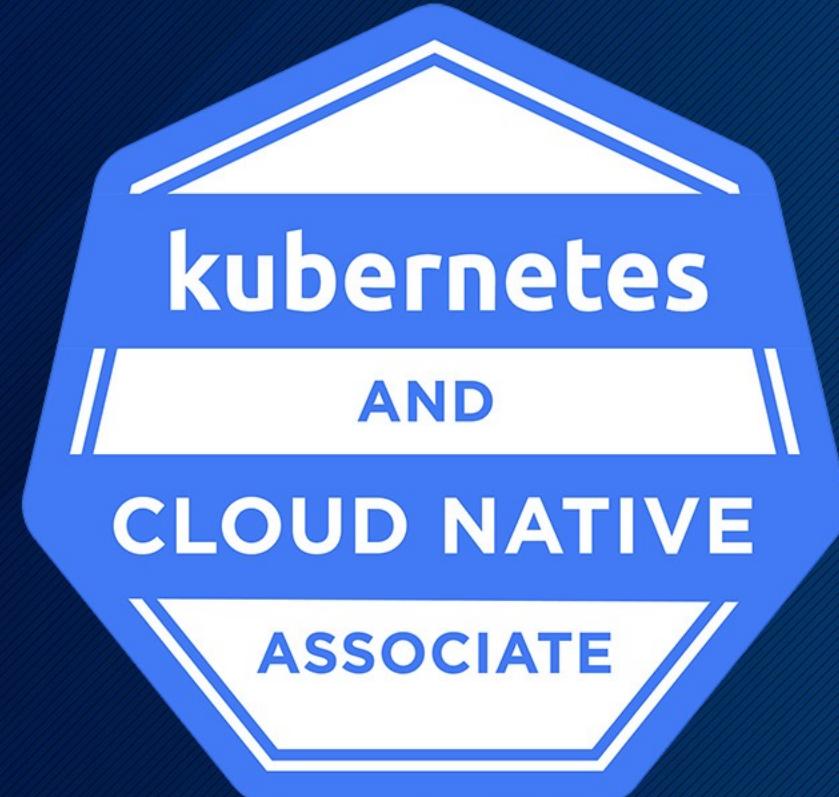


KCNA - Kubernetes and Cloud Native Associate

By Riyaz Sayyad



RIZMAXed



START COURSE »



Disclaimer: This content is copyrighted and strictly for personal use only.

- This content is reserved for people enrolled into the KCNA - Kubernetes and Cloud Native Associate course by Riyaz Sayyad. It is intended for personal use and exam preparation only, please do not share this content.
- All trademarks, logos and brand names are the property of their respective owners.



RIZMAXed

KCNA - Kubernetes and Cloud Native Associate

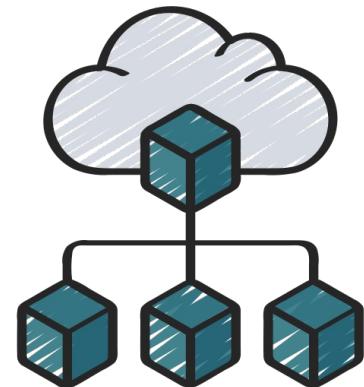
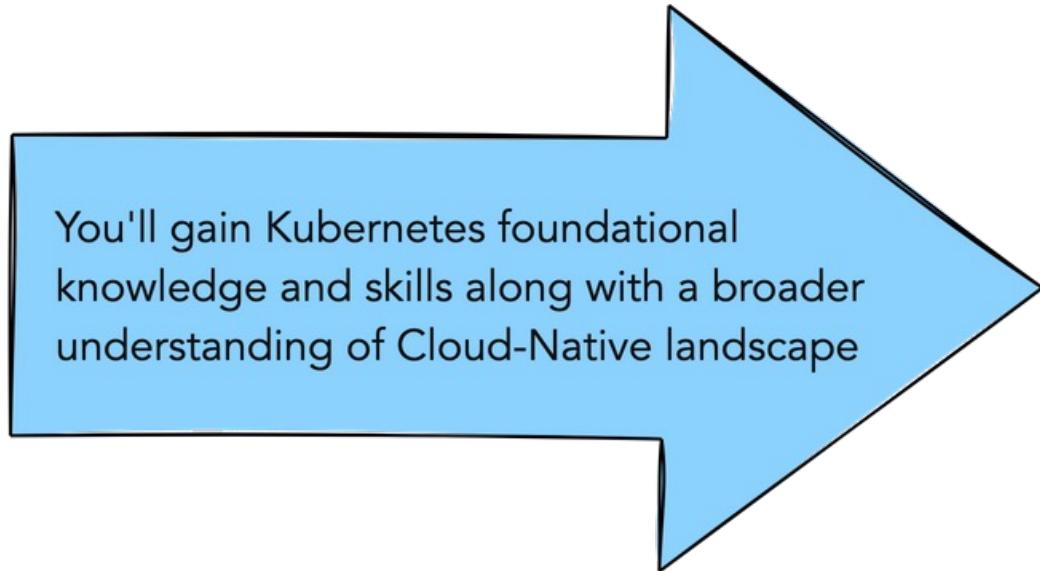
Hands-on Guide



Welcome!

RIZMAXed

- Let's get ready for the KCNA exam



- Recommended IT experience
 - Basic Linux Knowledge (Good to have)



About the exam

RIZMAXed

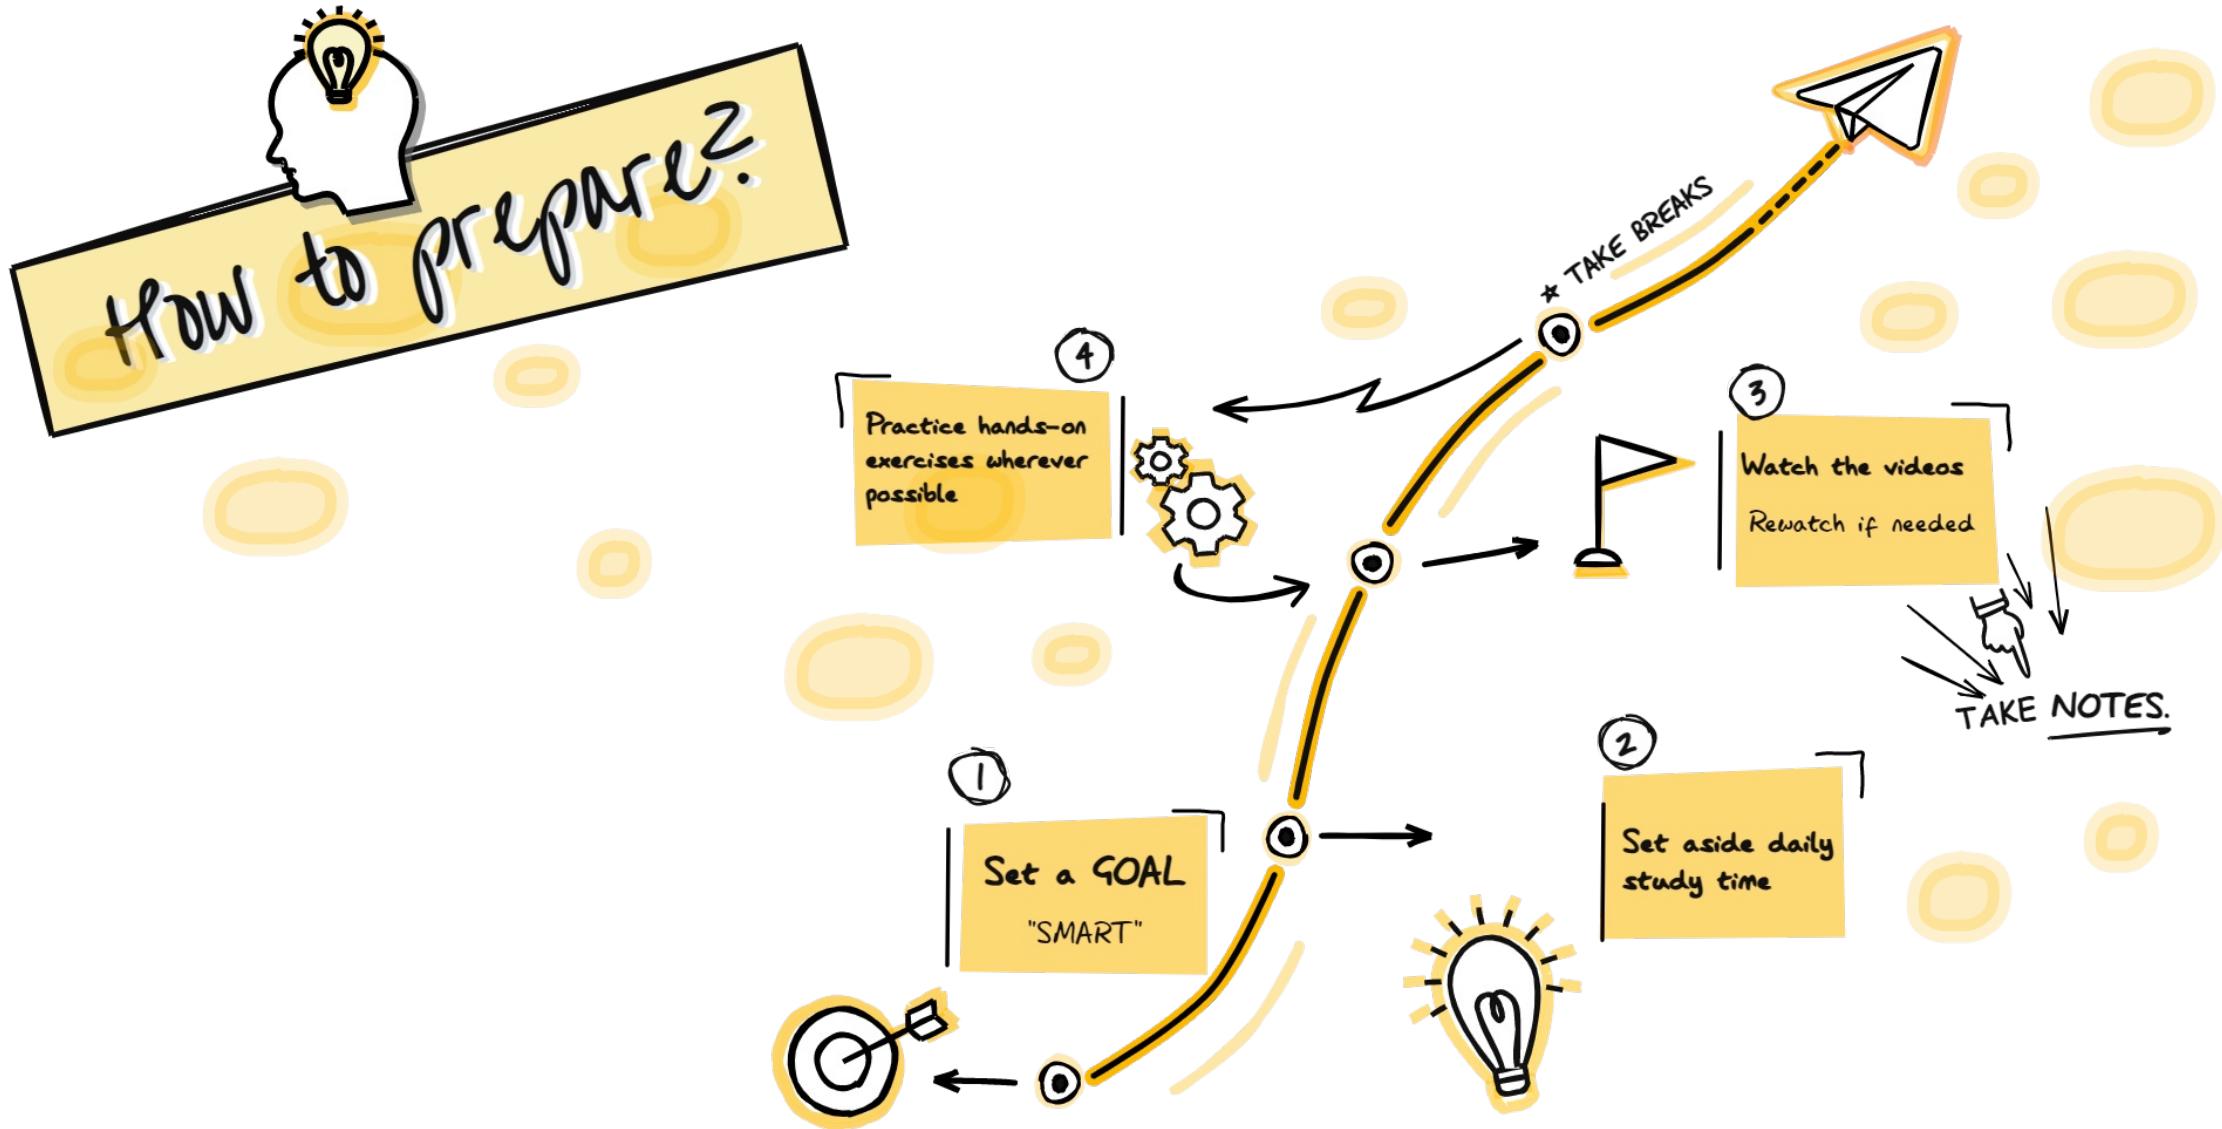
- Exam Overview
 - 90 minutes
 - 60 questions (MCQs / MRQs)
 - Passing score - 75%
 - No negative marking
 - Free retake
- Key Focus Areas
 - Kubernetes Fundamentals
 - Container Orchestration
 - Cloud Native Architecture
 - Cloud Native Observability
 - Cloud Native Application Delivery





Put your best foot forward!

RIZMAXed



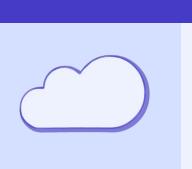


About me - I'm Riyaz!

RIZMAXed

- KCNA Certified
- 6x AWS Certified
- Over a decade of experience in AWS Cloud and Serverless
 - Built websites, apps, Cloud Native SaaS products
- Best-selling Instructor on Udemy
 - Serverless, Lambda, Databases
 - 100k+ students, 13k+ reviews
 - I love teaching! ❤️





Technology Overview

Containers, Kubernetes, and the cloud

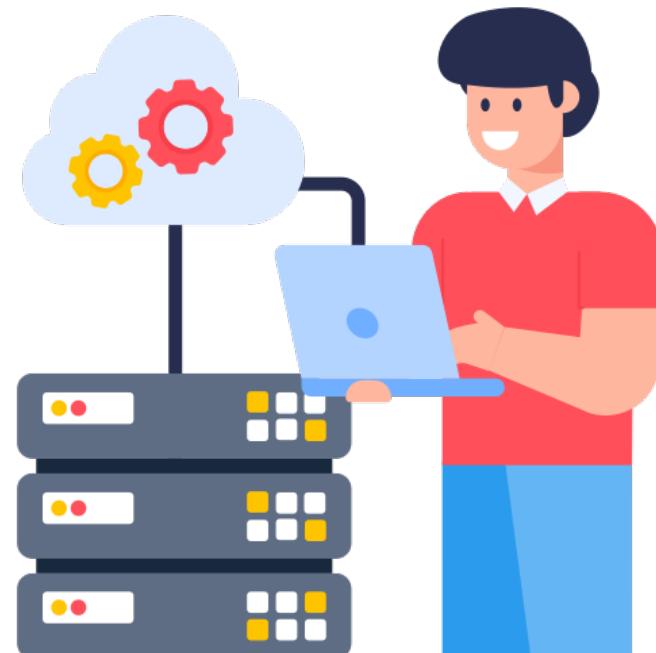




Cloud Computing

RIZMAXed

- On-demand, scalable IT resources
 - Compute, storage, databases, applications, etc.
- Pay-as-you-go (no upfront capital investment)
- Agility – deploy globally in minutes, faster TTM
- Main types of cloud computing – IaaS, PaaS, SaaS
- Cloud deployment models – Public, Private, Hybrid
- Public cloud – AWS, Azure, GCP
- Private cloud
 - Not same as the typical on-premise data center
 - On-premises data center w/ virtualization and resource management tools
- Hybrid = public + private cloud

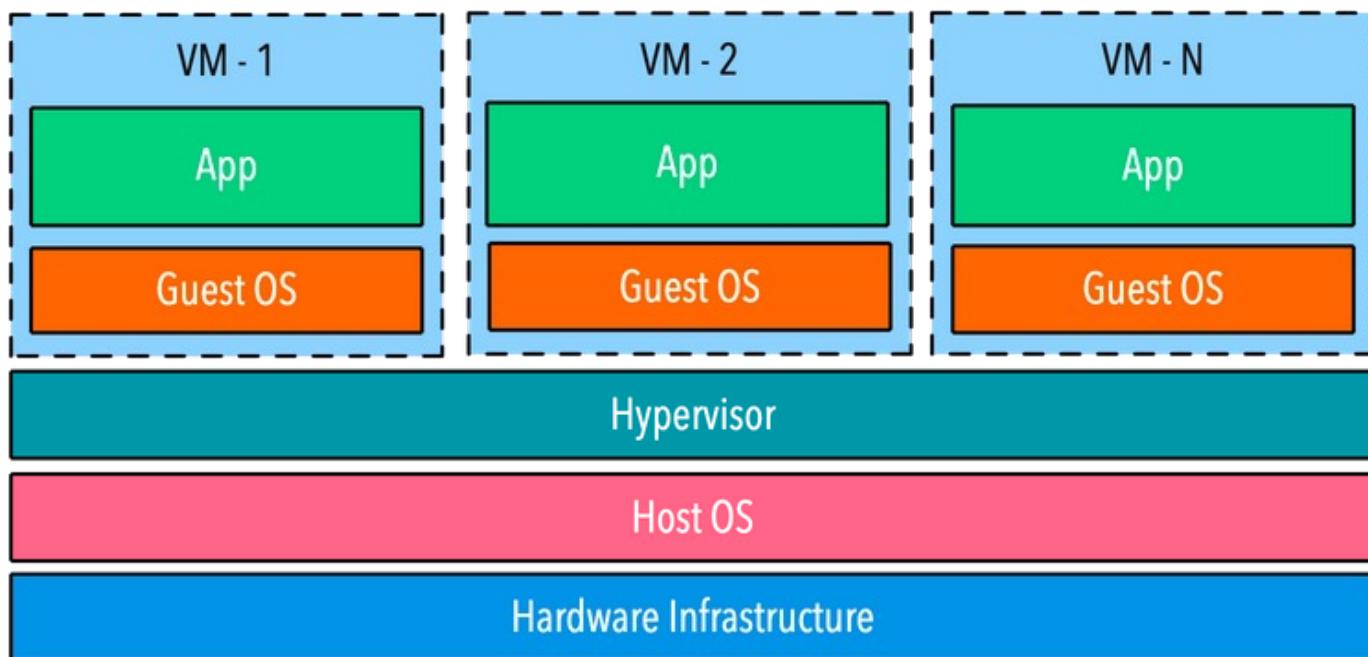




Virtualization

RIZMAXed

- running multiple operating systems on a single physical machine
- each operating system runs in a completely isolated environment
- managed or orchestrated by a hypervisor

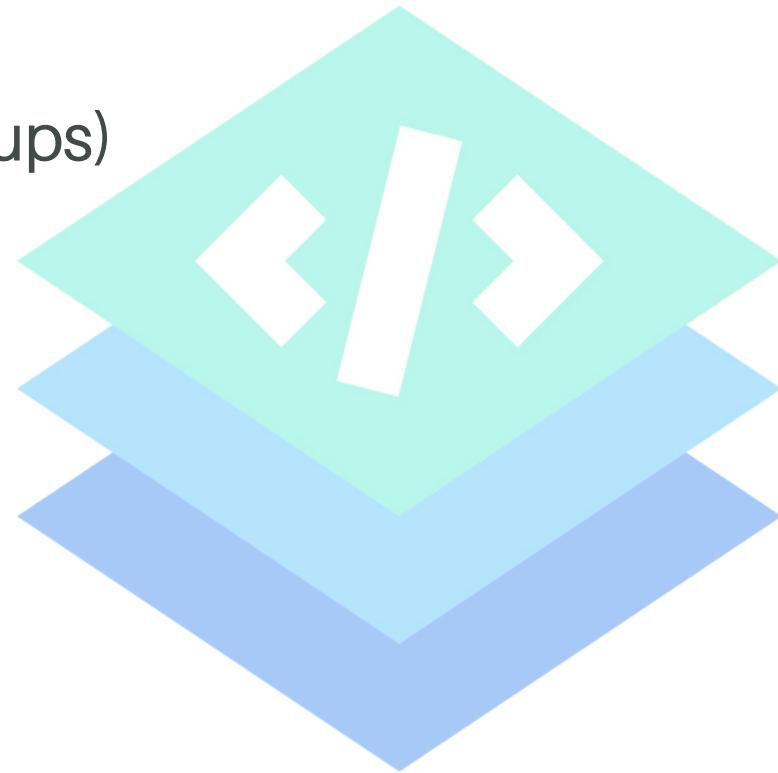
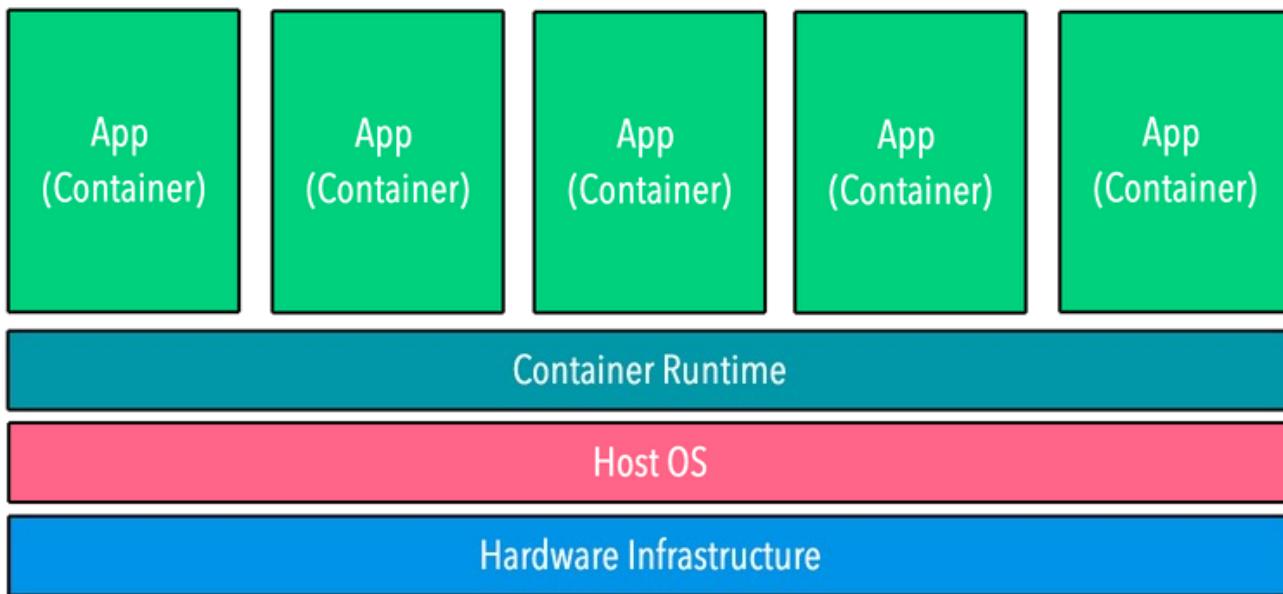




Containerization

RIZMAXed

- container image is a snapshot of an application + dependencies
- container is a running instance of a container image, each runs as an isolated process
- are portable, smaller in size, offer more efficient resource utilization
- multiple applications can run on fewer virtual machines
- isolation provided by Linux control groups feature (cgroups)





- provides a set of tools for developing, shipping, and running containerized applications
- a Docker image is defined using a **Dockerfile**
- **Docker Hub** is a central repository where users can store, share, and distribute Docker images
- container runtimes created by Docker
 - **runc** (lightweight, and portable)
 - **containerd** (industry-standard, simple, robust, and portable)
- runc was donated to OCI to help establish standardization
 - **OCI = Open Container Initiative** 
- containerd was donated to CNCF
 - **CNCF = Cloud Native Computing Foundation** 

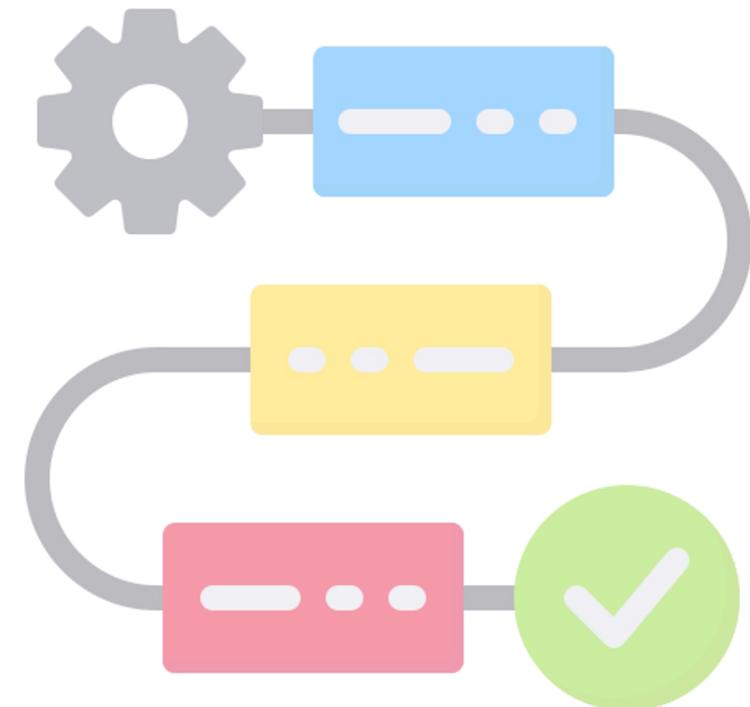
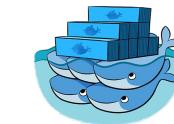




Container Orchestration

RIZMAXed

- automated management of containerized applications
 - including deployment, scaling, networking, and availability
- streamlines deploying and managing containers at scale
 - improves efficiency, reliability, and scalability
- Orchestration Platforms
 - Kubernetes
 - Docker Swarm
 - Apache Mesos
 - Amazon ECS
- Managed Kubernetes
 - EKS
 - AKS
 - GKE

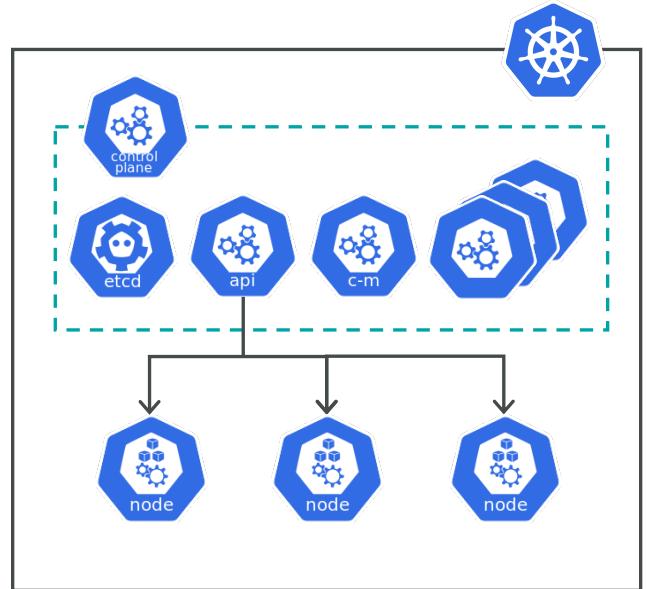




Kubernetes (K8s)

RIZMAXed

- open-source container orchestration platform
- developed by Google and donated to CNCF in 2014
- incubating 2016, graduated 2018
- Key features
 - automates the container scheduling and management
 - automatic scaling and self-healing
 - service discovery and load balancing
 - rolling updates and rollbacks
- master node controls and manages the cluster (control-plane)
- worker nodes host the containers and run the applications
- etcd acts as a distributed key value store to store cluster config data

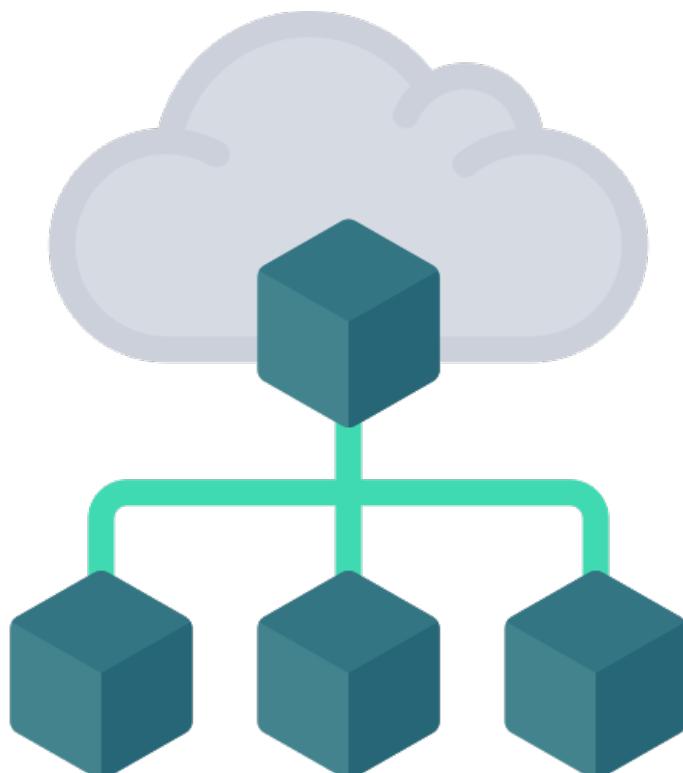




Cloud Native Computing

RIZMAXed

- = design and deployment of applications that fully leverage the benefits of cloud computing models (public, private, or hybrid)
- => scalable, resilient, agile, cost-effective, portable applications
- Key guiding principles
 - containerization and orchestration
 - microservices architecture
 - continuous integration/delivery/deployment
 - DevOps, GitOps, FinOps
- Use cases
 - distributed web apps
 - IoT and edge computing
 - big data and analytics





Environment Setup

RIZMAXed

- Install **Docker Desktop**
- Set up Multi-node Kubernetes cluster using **Minikube**
- Install **VS Code** (optional)



JAVASCRIPT OBJECT NOTATION

- Simple and popular format
- Used to exchange data over the APIs

```
{} cfn-template.json > ...
1   [
2     "Metadata": {
3       "AWS::CloudFormation::Interface": {
4         ...
5       }
6     },
7     "Parameters": {
8       ...
9     },
10    "Conditions": {
11      ...
12    },
13    "Mappings": {
14      ...
15    },
16    "Resources": {
17      "HANAIAMProfile": {
18        ...
19      },
20      "HanaDataVolume": {
21        ...
22      },
23      "HanaLogVolume": {
24        ...
25      },
26      "HanaOthersVolume": {
27        ...
28      },
29      "HanaSingleInstance": {
30        "Type": "AWS::EC2::Instance",
31        "Properties": {
32          ...
33        }
34      },
35      "AutoRecoverAlarm": {
36        ...
37      }
38    }
39  }
```



Hands-on Demo

RIZMAXed





YAML AIN'T MARKUP LANGUAGE

- Data serialization standard
- Human readable and computationally powerful
- Simple and popular text format
- Commonly used for writing templates and configuration files
- Do not use TAB characters for indentation. Use series of spaces instead. Modern IDEs like VS Code take care of this.

```
! cfn-template.yaml
  1   Metadata:
  2     'AWS::CloudFormation::Interface':
  3       ParameterGroups: ...
  69      ParameterLabels: ...
170     Parameters: ...
492    Conditions: ...
630    Mappings: ...
704   Resources:
705     HANAIAMProfile: ...
715     HanaDataVolume: ...
757     HanaLogVolume: ...
836     HanaOthersVolume: ...
880     HanaSingleInstance:
881       Type: 'AWS::EC2::Instance'
882       Properties:
883         BlockDeviceMappings: ...
896         IamInstanceProfile: !Ref HANAIAMProfile
897         ImageId: !Ref HANAAMIID
898         InstanceType: !Ref HANAINstanceType
899         KeyName: !Ref KeyName
900         Monitoring: true
901       NetworkInterfaces: ...
904       Tags: ...
927       UserData: !Base64 ...
1296      AutoRecoverAlarm: ...
```



Hands-on Demo

RIZMAXed





Linux Commands 101

RIZMAXed

- Quick refresher on Linux commands
 - we'll spin up Ubuntu container using Docker
- We'll cover
 - Basic navigation with Linux CLI
 - File and directory manipulation
 - Search and replace operations
 - File viewing and editing
 - Other common commands





Working with Containers

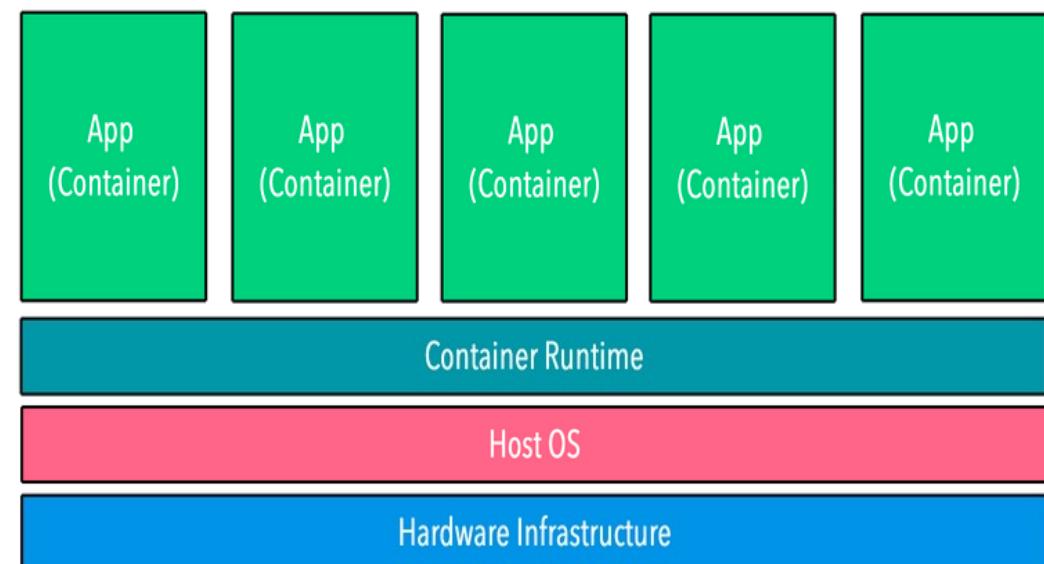
Docker fundamentals



Container Basics

RIZMAXed

- = packaged applications
 - application + dependencies + libs + config
 - consistent behavior across environments
 - lightweight, standalone, and portable
 - containers share the kernel of the OS
- Why use them?
 - Consistency, Isolation, Scalability
- Docker
 - Founded in 2010 with name dotCloud
 - Renamed and open-sourced in 2013
 - Most popular container platform
 - Easily create, deploy, manage containers

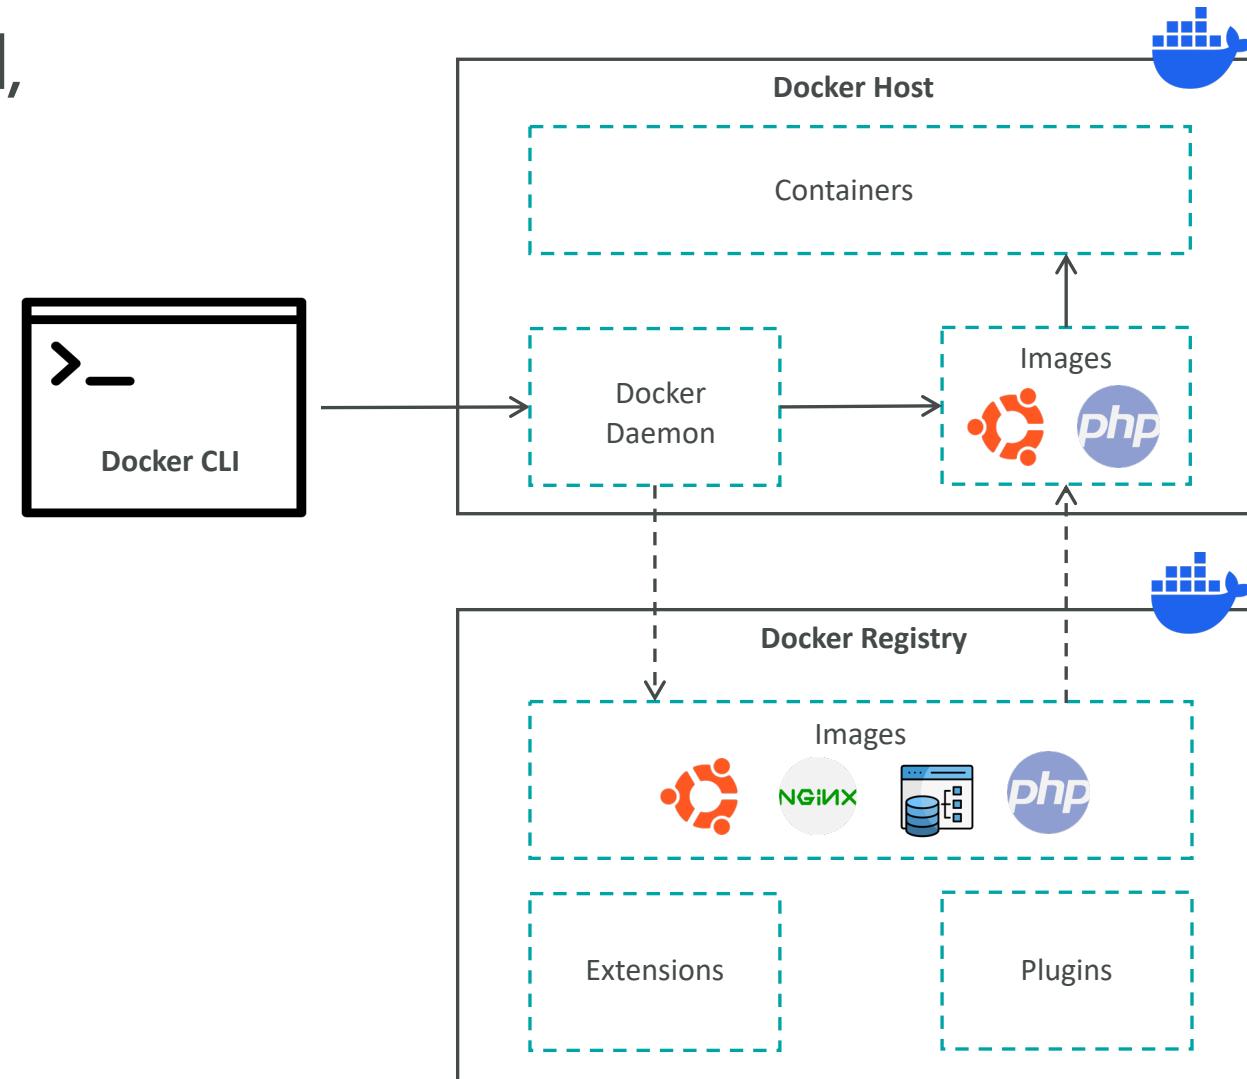




How Docker Works

RIZMAXed

- Docker Engine helps you build, ship, and run containers
- Key components
 - Docker Daemon
 - runs in background
 - Docker CLI
 - allows you to interact with Docker
 - Docker Registry
 - repository for Docker images
 - aka Docker Hub





Hands-on Demo - Docker Basics

RIZMAXed





Docker Commands Cheat Sheet

RIZMAXed

- **docker run <image>**
 - Create and start a new container from an image.
- **docker ps**
 - List all running containers.
- **docker ps -a**
 - List all containers (including stopped ones).
- **docker stop <container>**
 - Stop a running container.
- **docker start <container>**
 - Start a stopped container.
- **docker restart <container>**
 - Restart a container.
- **docker rm <container>**
 - Remove a stopped container.
- **docker images**
 - List all available images.
- **docker rmi <image>**
 - Remove an image.
- **docker pull <image>**
 - Download an image from a repository.
- **docker exec -it <container> <command>**
 - Run a command inside a running container.
- **docker logs <container>**
 - View logs from a container.
- **docker inspect <container>**
 - Display detailed information about a container.
- **docker network ls**
 - List all Docker networks.
- **docker volume ls**
 - List all Docker volumes.
- **docker build -t <tag> <path>**
 - Build an image from a Dockerfile.
- **docker login**
 - Log in to a Docker registry.
- **docker push <image>**
 - Push an image to a registry.
- **docker system info**
 - Display system-wide information about the Docker installation
- **docker <command> --help**
 - Look up help information on any docker command



Docker Port Mapping

RIZMAXed

- Publish all exposed ports to random ports

```
[riyaz@RizmaxMBP ~ % docker run --rm -d -P nginx:latest  
73c5df201d68c8ac99421b3672121d01be49f3af02910bf5813d6ef6304eba7b  
[riyaz@RizmaxMBP ~ %  
[riyaz@RizmaxMBP ~ % docker ps  
CONTAINER ID IMAGE COMMAND CREATED STATUS PORTS NAMES  
73c5df201d68 nginx:latest "/docker-entrypoint..." 6 seconds ago Up 5 seconds 0.0.0.0:55000->80/tcp kind_wiles  
riyaz@RizmaxMBP ~ %
```

- Publish container port(s) to specific port(s) on host

```
[riyaz@RizmaxMBP ~ % docker run --rm -d -p 8888:80 nginx:latest  
16bb0d90b84c1e0e507fc224353f52abe76c063541d8a0d41d74dd1663661d0d  
riyaz@RizmaxMBP ~ %  
[riyaz@RizmaxMBP ~ % docker ps  
CONTAINER ID IMAGE COMMAND CREATED STATUS PORTS NAMES  
16bb0d90b84c nginx:latest "/docker-entrypoint..." 39 seconds ago Up 39 seconds 0.0.0.0:8888->80/tcp charming_gr  
[riyaz@RizmaxMBP ~ %  
riyaz@RizmaxMBP ~ %
```



Docker Volume Binding

RIZMAXed

- Mount a file or directory on the host machine into a container

```
[riyaz@RizmaxMBP static-website % docker run --rm -d -p 8888:80 -v /Users/riyaz/sandbox/projects/static-website:/usr/share/nginx/html nginx:latest  
413e11b46e6e3a327824ab13644d2f85bed5877dafbc2063ab0406e938a1d4c1  
[riyaz@RizmaxMBP static-website % docker ps  
CONTAINER ID IMAGE COMMAND CREATED STATUS PORTS NAMES  
413e11b46e6e nginx:latest "/docker-entrypoint...." 12 seconds ago Up 11 seconds 0.0.0.0:8888->80/tcp competent_franklin  
riyaz@RizmaxMBP static-website %
```

Source Volume

:

Target Volume



Create a Container Image using Dockerfile

RIZMAXed

```
🐳 Dockerfile ×  
🐳 Dockerfile > ...  
1  FROM nginx:latest  
2  COPY index.html /usr/share/nginx/html  
3  EXPOSE 80  
4  CMD [ "nginx", "-g", "daemon off;" ]  
5  
6  
7  
8
```

```
$ docker build -t myapp .
```



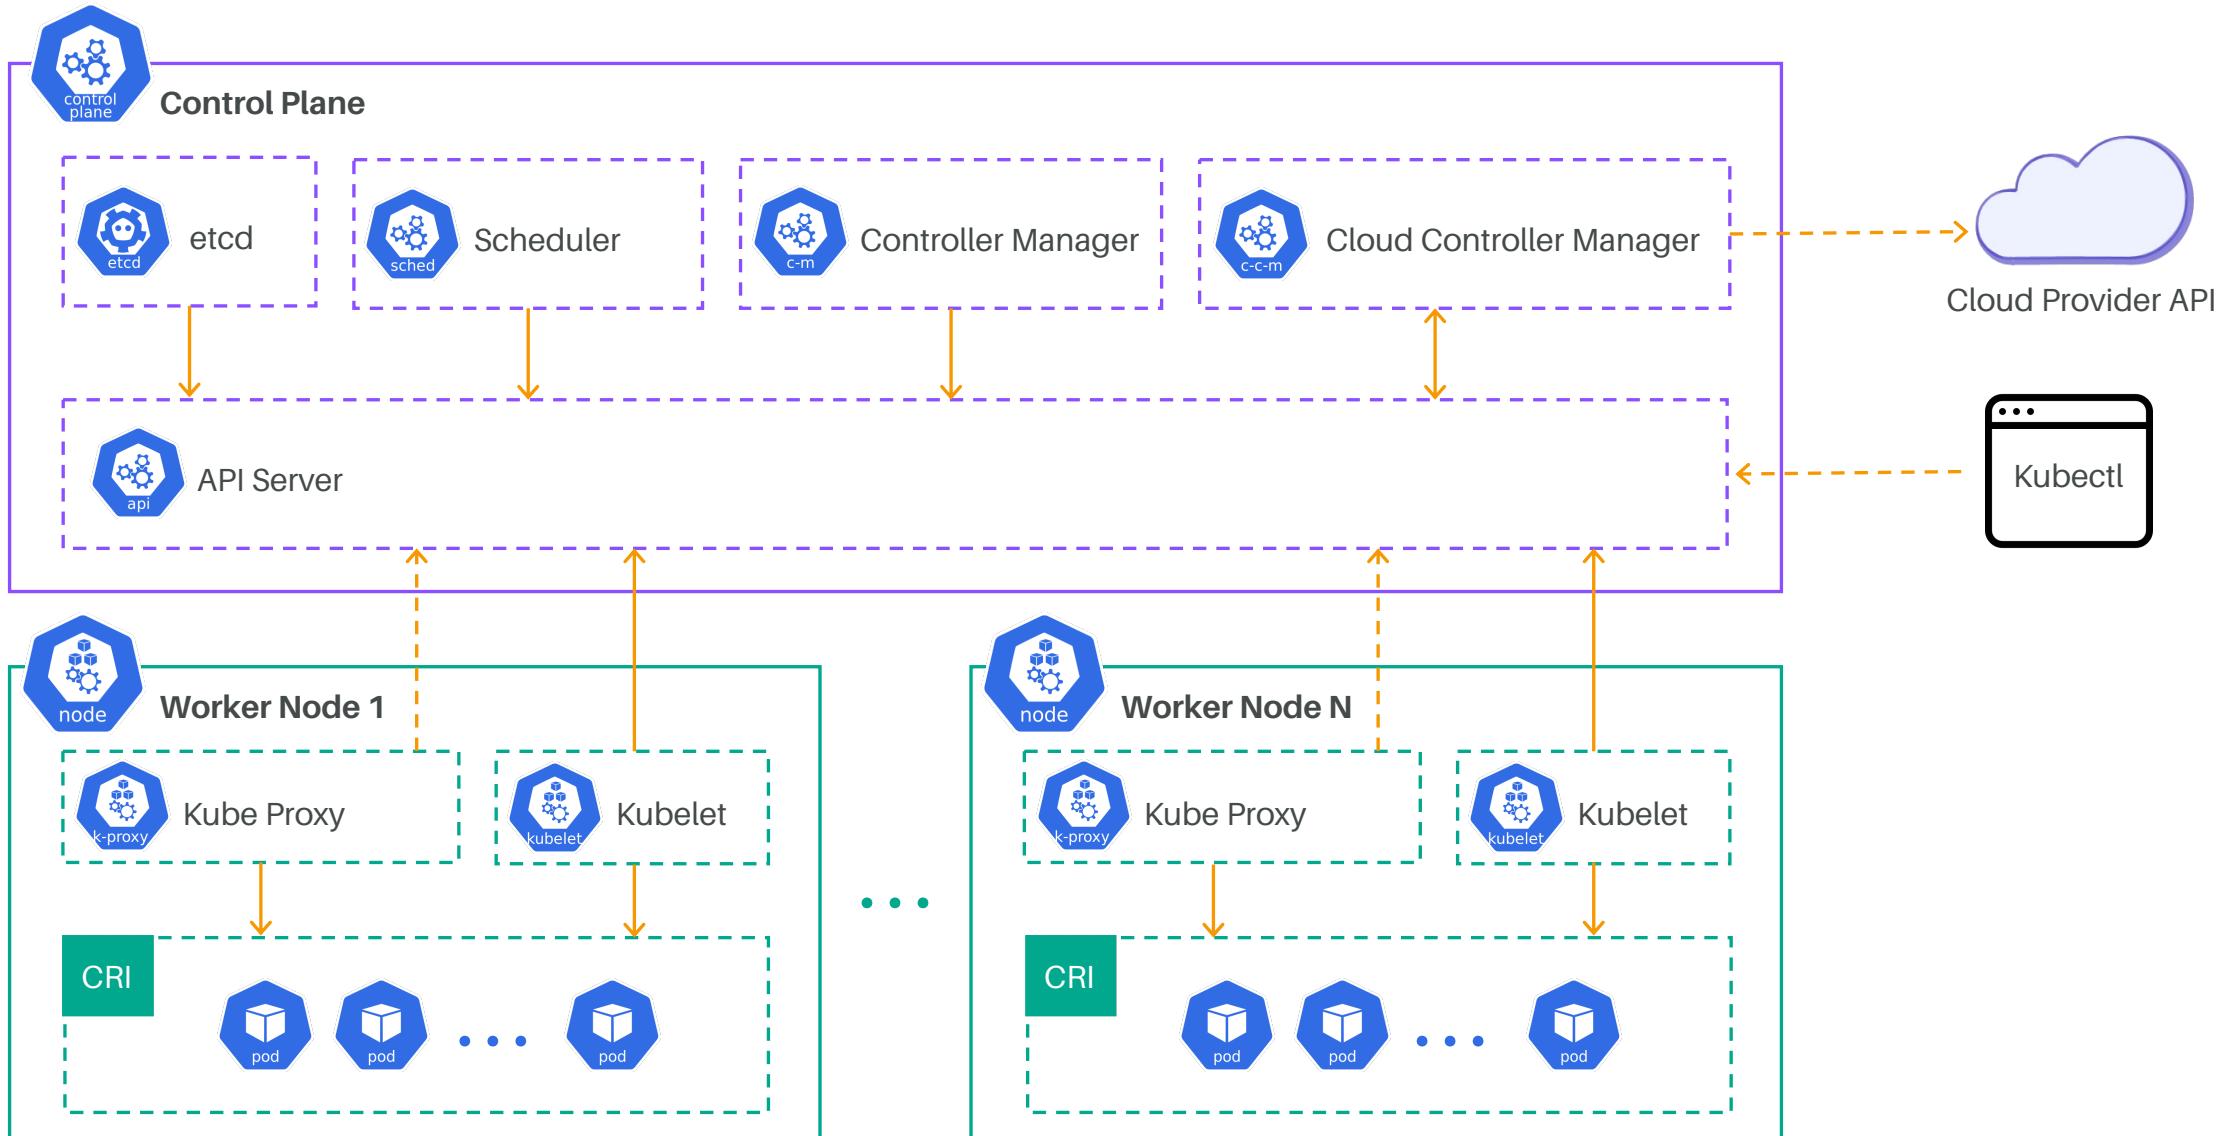
Kubernetes Overview

Get started with Kubernetes



Kubernetes Architecture

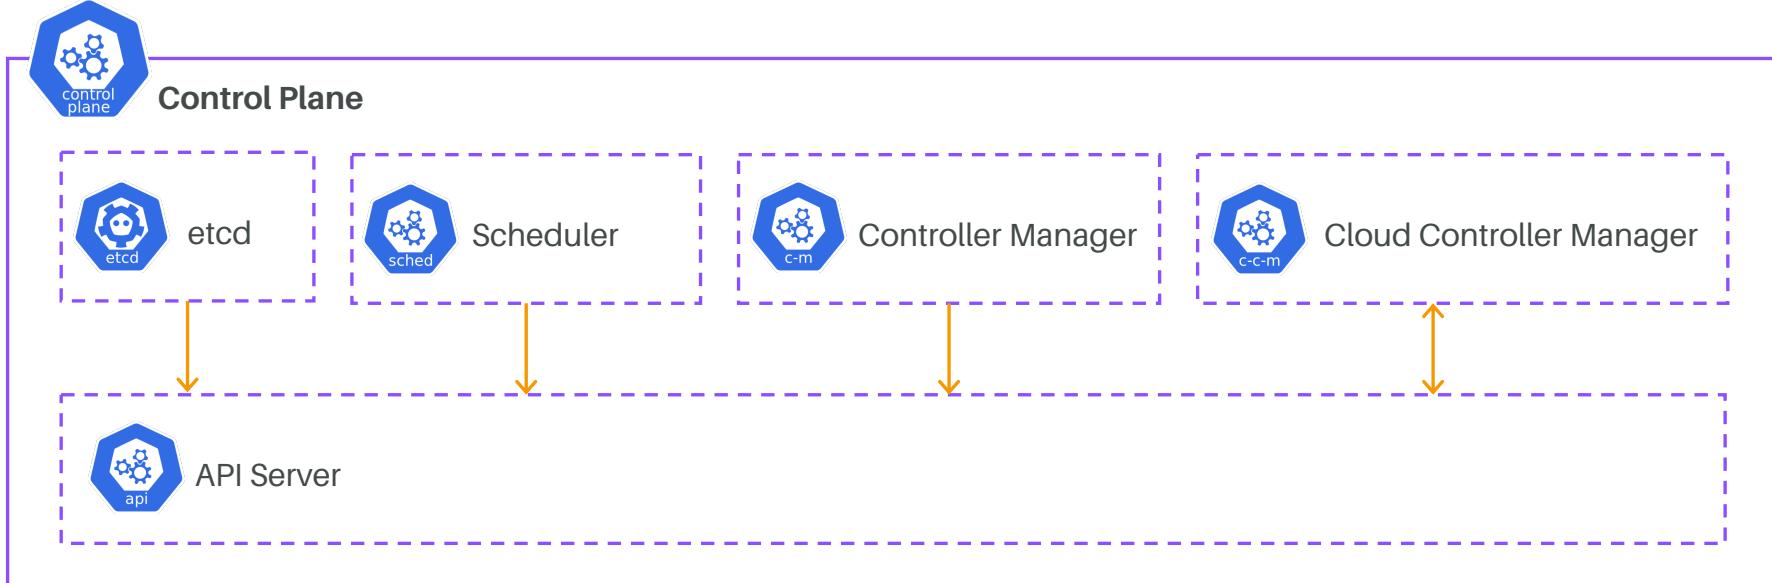
RIZMAXed





Control Plane

RIZMAXed



- Brain of the k8s cluster
- Manages and maintains the desired state of the cluster
 - resource management, scheduling, scaling, high-availability
- One or more control plane nodes
 - Single Control Plane Node - smaller or dev clusters
 - Highly Available Control Plane - production-grade clusters



Kube API Server

RIZMAXed



API Server

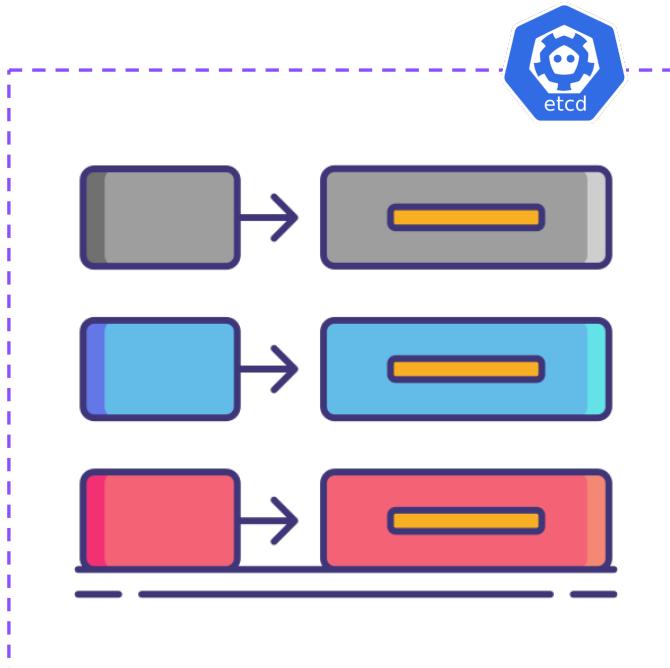
- Primary interface for interacting with a k8s cluster
- RESTful API – supports HTTP CRUD operations
- Resources are organized into API Groups
 - /api/v1 – core group **apiVersion: v1**
 - /apis/\$GROUP_NAME/\$VERSION- named groups **apiVersion: apps/v1**
- Custom resources can be defined using CRD
 - CRD = **Custom Resource Definition**
- Authentication – tokens, certificates, etc.
- Authorization – RBAC, ABAC, etc.
- Reconciles the desired state defined in the YAML or JSON manifests



etcd

RIZMAXed

- Distributed key-value store
- Primary datastore to store all cluster data
 - config data, state, and metadata of k8s objects
 - stores ConfigMaps and Secrets
- Reliable and consistent
 - highly available and fault-tolerant
 - implements Raft consensus algorithm
- Offers TLS encryption and authentication features
- Typically placed on control plane nodes
 - one etcd instance on each control plane node
- Can be placed on dedicated etcd cluster
 - multiple etcd instances in the etcd cluster

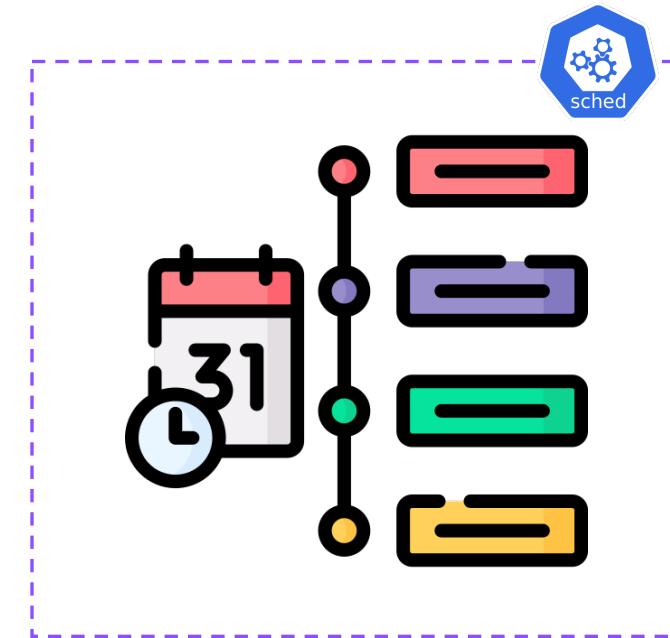




Kube Scheduler

RIZMAXed

- Scheduling
 - process to determine the best node to run a pod
- Kube Scheduler assigns pods to nodes
 - ensuring optimal resource utilization and
 - meeting application requirements
- 3-step operation
 - filtering – finds suitable nodes as per resource requirements
 - scoring – ranks found nodes
 - assignment – assigns pod to the highest-ranking node
- Pod placement can be controlled using
 - node name, node selectors, affinities, and taints/tolerations

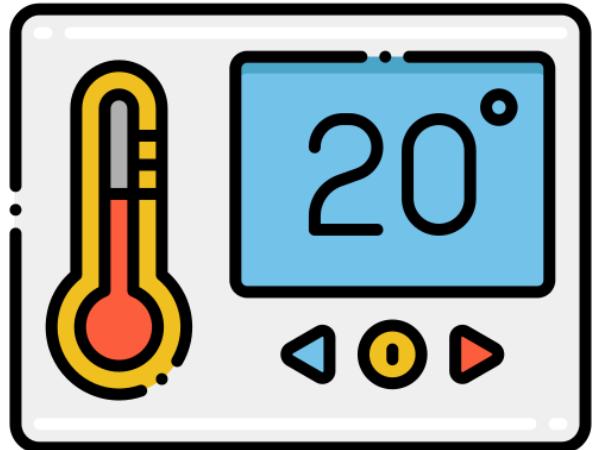




Controller Manager

RIZMAXed

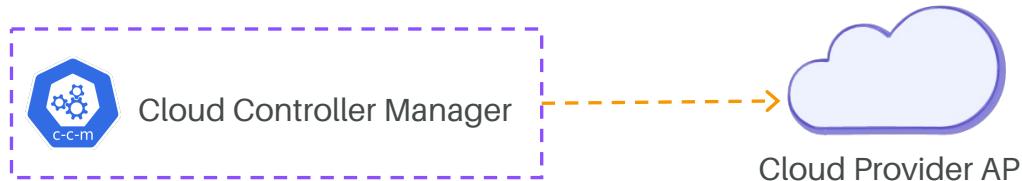
- Controllers
 - implement control loops (like a thermostat in a room)
 - continuously watch the cluster state via the API server
 - act to reconcile the current state with the desired state
- Controller Manager operates different controllers
 - Deployment Controller
 - Job controller etc.
- Can define custom controllers





Cloud Controller Manager (CCM)

RIZMAXed



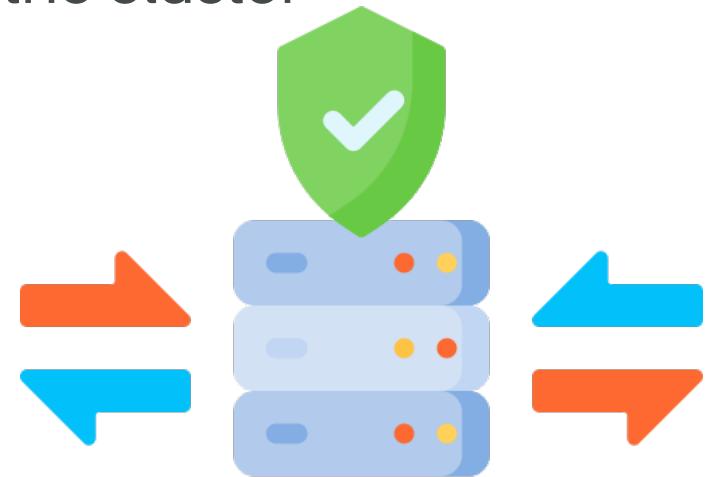
- bridge between Kubernetes and cloud providers
 - AWS, Azure, GCP, and more
- Abstracts cloud provider APIs
 - to manage resources like load balancers, volumes, instances
 - allows any cloud provider to integrate with k8s
- Key functions
 - Resource provisioning - persistent volumes, load balancers etc.
 - Node management - create/delete instances, auto-scaling
 - Networking - communication between pods and external services



Kube Proxy

RIZMAXed

- Runs on all worker nodes
- Acts as a Network Proxy
 - managing network connectivity to k8s services within the cluster
- Abstracts the internal network details
 - providing a consistent view of services to applications
- Enforces network policies
 - to control traffic between pods and services
- Implements load balancing for services
 - routing traffic to healthy pods
- Facilitates service discovery

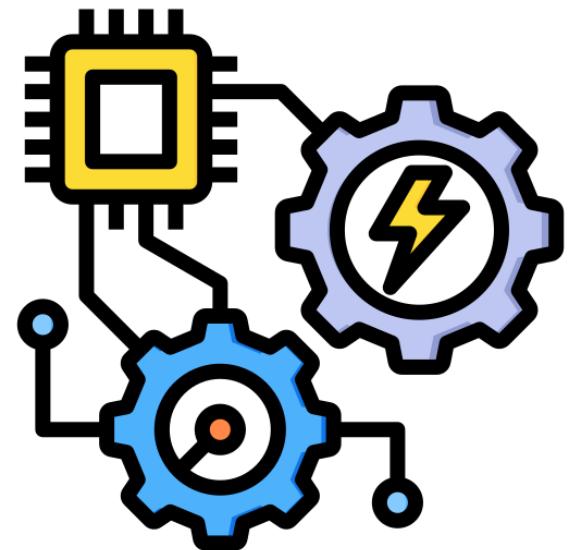




Kubelet

RIZMAXed

- Primary **node agent** that runs on each node
- Ensures pods are instantiated, running, and terminated
 - according to the desired state defined in pod specifications
- Handles resource allocation to pods
 - CPU, memory, and storage
- Monitors node and container health
 - uses liveness probes to check container health
 - e.g. HTTP Probe, Exec Probe, TCP probe
 - restarts containers that fail the liveness probe
- Interfaces with container runtime to manage containers within pods





Container Runtime Interface (CRI)

RIZMAXed

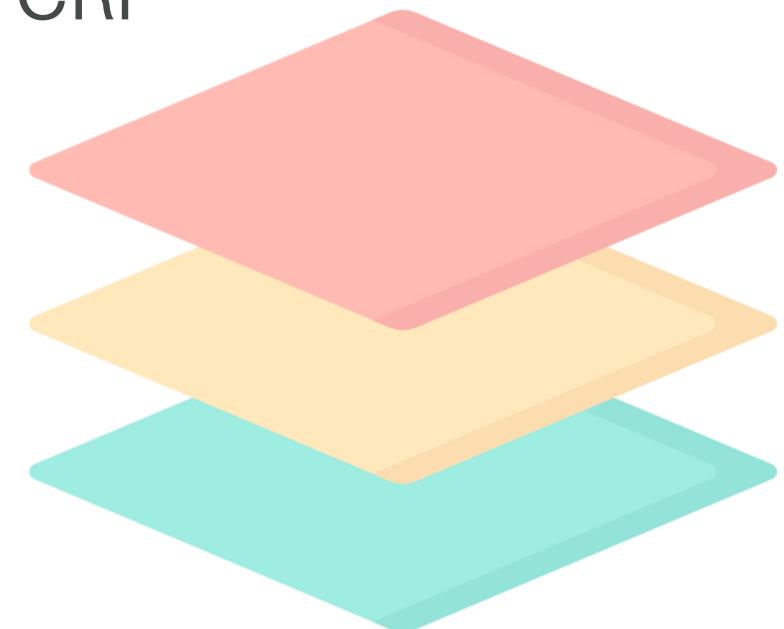
- Container runtime runs containers within pods
- CRI is an abstraction layer on top of Container Runtime
 - enables the kubelet to use a wide variety of container runtimes
 - => flexibility and interoperability across runtime environments
- Supports any container runtime that implements CRI
 - e.g. containerd, CRI-O, Kata Containers, Firecracker
- Docker (dockershim) is no longer supported



cri-o



Firecracker





Kubeconfig

RIZMAXed

- Configuration file
 - cluster info, auth credentials, and context settings
- Clusters
 - cluster endpoint, Server CA info
- Users
 - user authentication creds (tokens, certs)
- Contexts
 - defines cluster access configuration
 - cluster, user, namespace, + other settings
- Typically stored in
 - **~/.kube/config**
 - or in **KUBECONFIG** env var

```
Rizmax ~ $ cat ~/.kube/config
apiVersion: v1
clusters:
- cluster:
  certificate-authority: /Users/riyaz/.minikube/ca.crt
  extensions:
  - extension:
    last-update: Wed, 13 Mar 2024 07:03:53 EDT
    provider: minikube.sigs.k8s.io
    version: v1.32.0
    name: cluster_info
    server: https://127.0.0.1:59873
    name: minikube
contexts:
- context:
  cluster: minikube
  extensions:
  - extension:
    last-update: Wed, 13 Mar 2024 07:03:53 EDT
    provider: minikube.sigs.k8s.io
    version: v1.32.0
    name: context_info
    namespace: default
    user: minikube
    name: minikube
  current-context: minikube
  kind: Config
  preferences: {}
users:
- name: minikube
  user:
    client-certificate: /Users/riyaz/.minikube/profiles/minikube/client.crt
    client-key: /Users/riyaz/.minikube/profiles/minikube/client.key
```



Semantic Versioning (SemVer)

RIZMAXed

- Kubernetes follows **SemVer**
 - Kubernetes Releases
 - API versions
 - Resource Definitions
 - Component Versions
- Kubernetes releases follow **MAJOR.MINOR.PATCH** format
 - major releases may introduce breaking changes
 - 1.0.0, 2.0.0
 - minor releases introduce new features
 - in a backward-compatible manner
 - 1.19.0, 1.20.0
 - patch releases contain bug fixes and minor enhancements
 - 1.19.1, 1.19.2
- API versions may be labeled as alpha, beta, or stable



Kubernetes Resources

RIZMAXed

- = API objects in k8s
 - Pod, Service, Deployment, ReplicaSet, etc.
- Scoped by **namespaces**
 - divide cluster resources into virtual clusters
 - isolate resources of different projects / teams / environments within a cluster
 - can set resource quotas at namespace level
 - CPU, memory, etc.
 - if no namespace is specified, **default** namespace is used
- Defined using YAML or JSON manifests
 - => **Declarative, IaC**
- Can also be created using CLI commands
 - => **Imperative**
- Use CRD to define custom resources
 - CRD = **Custom Resource Definition**

```
$ kubectl api-resources
```

Rizmax ~ \$ kubectl api-resources	NAME	SHORTNAMES	APIVERSION	NAMESPACED	KIND
	bindings		v1	true	Binding
	componentstatuses	cs	v1	false	ComponentStatus
	configmaps	cm	v1	true	ConfigMap
	endpoints	ep	v1	true	Endpoints
	events	ev	v1	true	Event
	limitranges	limits	v1	true	LimitRange
	namespaces	ns	v1	false	Namespace
	nodes	no	v1	false	Node
	persistentvolumeclaims	pvc	v1	true	PersistentVolumeClaim
	persistentvolumes	pv	v1	false	PersistentVolume
	pods	po	v1	true	Pod
	podtemplates		v1	true	PodTemplate
	replicationcontrollers	rc	v1	true	ReplicationController
	resourcequotas	quota	v1	true	ResourceQuota
	secrets		v1	true	Secret
	serviceaccounts	sa	v1	true	ServiceAccount
	services	svc	v1	true	Service
	mutatingwebhookconfigurations		admissionreg	false	MutatingWebhookConfiguration
	validatingwebhookconfigurations		admissionreg	false	ValidatingWebhookConfiguration
	customresourcedefinitions	crd, crds	apiextension	false	CustomResourceDefinition
	apiservices		apiregistrat	false	APIService
	controllerrevisions		apps/v1	true	ControllerRevision
	daemonsets	ds	apps/v1	true	DaemonSet
	deployments	deploy	apps/v1	true	Deployment
	replicasets	rs	apps/v1	true	ReplicaSet
	statefulsets	sts	apps/v1	true	StatefulSet



Hands-on Demo

RIZMAXed





Imperative vs Declarative

RIZMAXed

- Imperative approach
 - specify exact sequence of steps to provision resources
 - give explicit instructions for implementation
 - e.g. creating resources with CLI commands

```
$ kubectl run mypod --image=nginx:stable
```
- Declarative approach – IaC, preferred
 - specify desired end state
 - define desired state using YAML manifests
 - k8s reconciles current state with desired state
 - required top-level attributes💡
 - apiVersion
 - kind
 - metadata
 - spec

```
! first-pod.yaml x
! first-pod.yaml > {} spec > [ ] containers > {} 0 > [ ] ports > {} 0
          io.k8s.api.core.v1.Pod (v1@pod.json)
1   apiVersion: v1
2   kind: Pod
3   metadata:
4     name: mypod
5   labels:
6     name: mypod
7   spec:
8     containers:
9       - name: nginx
10      image: nginx:stable
11      resources:
12        limits:
13          memory: "128Mi"
14          cpu: "500m"
15      ports:
16        - containerPort: 80
17
```



Pods

RIZMAXed

- The smallest deployable unit in k8s
- One pod runs one or more containers
 - typically, one app container => single-container pod
 - can have multiple app containers, rare
- Multi-container pods
 - app container + helper or sidecar containers
- All containers in a pod share the same
 - IP address, hostname, storage volumes
- Pod Lifecycle
 - Pending > Running > Succeeded
 - Failed, Unknown
- Restart Policy
 - kubelet is designed to restart exited containers
 - **restartPolicy: Always | OnFailure | Never**
 - default restartPolicy = Always
 - can be defined at pod-level or container-level

```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: mypod
5    labels:
6      name: mypod
7  spec:
8    containers:
9      - name: nginx
10        image: nginx:stable
11        resources:
12          limits:
13            memory: "128Mi"
14            cpu: "500m"
15        ports:
16          - containerPort: 80
17        restartPolicy: OnFailure
18
```



Labels and Annotations

RIZMAXed

- Labels

- key-value pairs attached to k8s objects
- for identification, grouping, and selection
- enable flexible querying, filtering, and organizing of resources
- used with selectors for resource targeting

- Annotations

- additional metadata attached to k8s objects
- for providing extra information
- for documentation, debugging, or tooling purposes
- limit the use of annotations to essential metadata to avoid cluttering

```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app: myapp
    environment: production
  annotations:
    description: "Frontend for production environment"
    buildVersion: "v1.2.3"
    imageregistry: "https://hub.docker.com/"
    kubernetes.io/change-cause: "Bug fix 871"
```



ReplicaSet

RIZMAXed

- Ensures that a specified # of identical pods are running at any given time
- => high availability and scalability
- Used for **stateless** apps
- Can scale down to zero
- selector is used to identify pods to run replicas
- Deployment is k8s object used to manage ReplicaSets
- Recommended to use Deployments
- Use ReplicaSets only if you need to manage pod replicas directly

```
1  apiVersion: apps/v1
2  kind: ReplicaSet
3  metadata:
4    name: myapp-replicaset
5  spec:
6    replicas: 3
7    selector:
8      matchLabels:
9        app: myapp
10   template:
11     metadata:
12       labels:
13         app: myapp
14     spec:
15       containers:
16         - name: myapp-container
17           image: myapp-image
18
```



Deployment

RIZMAXed

- A Deployment manages ReplicaSets
- declarative updates to pods
- => high availability and scalability
- + updates and rollbacks
- Update Strategies
 - RollingUpdate (default)
 - Recreate
 - use `maxUnavailable` and `maxSurge` params to control behavior in case of rolling updates
- Can implement canary deployments by creating multiple deployments
- Deployments are **stateless**

```
$ kubectl apply -f myapp-deployment.yaml
```

```
$ kubectl scale deployment/myapp-deployment --replicas=10
```

```
$ kubectl rollout undo deployment/myapp-deployment
```

```
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: myapp-deployment
5  spec:
6    replicas: 3
7    strategy:
8      type: RollingUpdate
9      rollingUpdate:
10        maxSurge: 1
11        maxUnavailable: 1
12    selector:
13      matchLabels:
14        app: myapp
15    template:
16      metadata:
17        labels:
18          app: myapp
19      spec:
20        containers:
21          - name: myapp-container
22            image: myapp-image
```



Init and Sidecar Containers

RIZMAXed

• Init Container

- runs before the app containers start
- to run any utilities or setup scripts
 - database schema initialization
 - data pre-processing
 - configuration preparation
- always run to completion
- must complete successfully before the next one starts

• Sidecar Container

- run alongside the main app container
- ignores the Pod-level restartPolicy
- to provide additional functionality
 - logging, monitoring, or proxying for the main app
 - e.g. Prometheus as a sidecar to collect metrics

```
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: myapp
5    labels:
6      app: myapp
7  spec:
8    replicas: 2
9    selector:
10      matchLabels:
11        app: myapp
12    template:
13      metadata:
14        labels:
15          app: myapp
16      spec:
17        containers:
18          - name: myapp-container
19            image: myapp-image
20        initContainers:
21          - name: init-setup
22            image: busybox
23            command: ["sh", "-c", "wget -O /config/config.yaml http://
24 config-server/config.yaml"]
25          - name: wait-for-db
26            image: busybox
27            command: ["sh", "-c", "until nc -z -v -w30 db-hostname 3306;
28 do echo 'Waiting for DB connection...'; sleep 5; done"]
29          - name: logshipper
30            image: busybox
31            restartPolicy: Always
32            command: ["sh", '-c', 'tail -F /opt/logs.txt']
```



StatefulSet

RIZMAXed

- Used to manage deployment and scaling of stateful apps
- Runs identical pods just like Deployment/ReplicaSets
- Maintains sticky identity for each pod
 - pod identifier(name/network hostname) remains the same between rescheduling
- used for **stateful** applications
- Ensures stable persistent storage for each pod in the set
- selector is used to identify pods to run replicas
- Currently requires a **Headless Service**
- Ideal for running databases, caching systems, messaging queues
- Update Strategies
 - RollingUpdate (default)
 - OnDelete

```
1  apiVersion: apps/v1
2  kind: StatefulSet
3  metadata:
4    name: my-statefulset
5  spec:
6    replicas: 3
7    serviceName: my-statefulset
8    selector:
9      matchLabels:
10     app: my-statefulset
11   template:
12     metadata:
13       labels:
14         app: my-statefulset
15   spec:
16     containers:
17       - name: my-container
18         image: my-image
19         ports:
20           - containerPort: 80
```



DaemonSet

RIZMAXed

- ensures that a specified pod runs on every node in a cluster
 - or on some nodes (using filters)
- Ideal for
 - system daemons
 - log collectors
 - monitoring agents
 - other tasks that should run on all nodes
- selector is used to identify pods to run daemons
- Update Strategies
 - RollingUpdate (default)
 - OnDelete

```
1  apiVersion: apps/v1
2  kind: DaemonSet
3  metadata:
4    name: logshipper
5    labels:
6      app: logshipper
7  spec:
8    selector:
9      matchLabels:
10     app: logshipper
11   template:
12     metadata:
13       labels:
14         app: logshipper
15     spec:
16       containers:
17         - name: logshipper
18           image: busybox
19           command: ["sh", "-c", "tail -F /opt/logs.txt"]
20
```



Job

RIZMAXed

- Used to execute one-time tasks that run to completion
- Creates one or more pods to execute tasks
- Retries execution of pods until a specified # of them successfully terminate
- for one-time tasks like batch processing, data processing, etc.
- Pod restartPolicy can only be Never or OnFailure
- Deleting a Job deletes the pods created

```
apiVersion: batch/v1
kind: Job
metadata:
  name: my-job
spec:
  template:
    spec:
      containers:
        - name: my-container
          image: my-image
          command: ["echo", "Hello from my job!"]
      restartPolicy: Never
      backoffLimit: 4
```



CronJob

RIZMAXed

- Creates Jobs on a predefined schedule
- Similar to Unix Cron jobs
- Schedule is defined using a crontab expression
- For automation of repetitive tasks, or scheduled tasks

```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: my-cronjob
spec:
  schedule: "0 22 * * 1-5"
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: my-container
              image: my-image
              imagePullPolicy: IfNotPresent
              command: ["echo", "Hello from my cronjob!"]
          restartPolicy: OnFailure
```



ConfigMaps and Secrets

RIZMAXed

- Store data as key-value pairs
- ConfigMaps store non-sensitive config data
- Secrets store sensitive data
 - not encrypted, only base64 encoded
 - can be encrypted with additional config
- Consumed by pods
- Can be mounted as data volumes or exposed as env vars
- Decouples config data from application code
- Can be immutable after creation

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: my-configmap
data:
  key1: value1
  key2: value2
immutable: true
```

```
apiVersion: v1
kind: Secret
metadata:
  name: my-secret
type: Opaque
data:
  username: <base64-encoded-value>
  password: <base64-encoded-value>
immutable: true
```



Hands-on Demo

RIZMAXed





Kubernetes Deep Dive

Container Orchestration



Scheduling

RIZMAXed

- Scheduling
 - process to determine the best node to run a pod
- Kube Scheduler assigns pods to nodes
 - ensuring optimal resource utilization and
 - meeting application requirements
- 3-step operation
 - filtering – finds suitable nodes as per resource requirements
 - scoring – ranks found nodes
 - assignment – assigns pod to the highest-ranking node
- Pod placement can be controlled using
 - node name, node selectors, affinities, and taints/tolerations

```
! first-pod.yaml x
! first-pod.yaml > {} spec > [ ] containers > {} 0 > [ ] ports > {} 0
          io.k8s.api.core.v1.Pod (v1@pod.json)
1   apiVersion: v1
2   kind: Pod
3   metadata:
4     name: mypod
5   labels:
6     name: mypod
7   spec:
8     containers:
9       - name: nginx
10      image: nginx:stable
11
12      resources:
13        limits:
14          memory: "128Mi"
15          cpu: "500m"
16
17      ports:
18        - containerPort: 80
```



Controlling Pod Placement

RIZMAXed

- Specify **nodeName** in your pod spec
- Specify node labels with **nodeSelector**

```
Rizmax ~ $ kubectl label node minikube-m02 env=dev
node/minikube-m02 labeled
```

```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: nginx
5  spec:
6    containers:
7      - name: nginx
8        image: nginx
9    nodeSelector:
10      env: dev
11
```

```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: nginx
5  spec:
6    containers:
7      - name: nginx
8        image: nginx
9    nodeName: worker-01
10
```



Controlling Pod Placement

RIZMAXed

- **Node Affinity and Anti-affinity**
 - similar to nodeSelector
 - provides more control node selection
 - use logical operators to select nodes
- **Node Affinity**
 - ensures that pods are scheduled on nodes with certain labels
 - Use operators like In, Exists, Gt, Lt
- **Node Anti-affinity**
 - prevents pods from being scheduled on nodes with certain labels
 - use operators like NotIn, DoesNotExist

Similar to nodeAffinity, use **podAffinity** to co-locate pods or **podAntiAffinity** to prevent them from being scheduled together.

```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: nginx
5  spec:
6    affinity:
7      nodeAffinity:
8        requiredDuringSchedulingIgnoredDuringExecution:
9          nodeSelectorTerms:
10            - matchExpressions:
11              - key: env
12                operator: In
13                  values:
14                    - dev
15                    - test
16              - key: env
17                operator: NotIn
18                  values:
19                    - prod
20      containers:
21        - name: nginx
22          image: nginx
```

requiredDuringSchedulingIgnoredDuringExecution
=> enforce the rule

preferredDuringSchedulingIgnoredDuringExecution
=> suggest using the rule



Controlling Pod Placement

RIZMAXed

- Taints and tolerations
 - ensure that pods are not scheduled onto inappropriate node
- Taints
 - labels applied to nodes that repel certain pods
 - `kubectl taint nodes <node-name> <key>=<value>:<effect>`
 - effect: NoExecute | NoSchedule | PreferNoSchedule
- Tolerations
 - applied to pods
 - allow pods to tolerate node taints and still be scheduled on tainted nodes

```
Rizmax ~ $ kubectl taint nodes minikube-m02 team=noram:NoSchedule  
node/minikube-m02 tainted
```

```
1  apiVersion: v1  
2  kind: Pod  
3  metadata:  
4    name: nginx  
5    labels:  
6      env: test  
7  spec:  
8    containers:  
9      - name: nginx  
10        image: nginx  
11    tolerations:  
12      - key: "team"  
13        operator: "Equal"  
14        value: "noram"  
15        effect: "NoSchedule"  
16
```



Static Pods

RIZMAXed

- = pods managed directly by kubelet on a node
 - not by the Kubernetes API server
 - managed outside the control plane (local management)
- Suitable for lightweight, node-specific tasks
- Simply place the pod manifest YAML or JSON files inside
/etc/kubernetes/manifests
- Kubelet automatically detects and manages static pods
- Lack features like scaling, rolling updates, or advanced scheduling options
- Use judiciously considering trade-offs in manageability and flexibility



Pod Disruption Budgets (PDB)

RIZMAXed

- Ensure pod availability during voluntary disruptions
 - i.e. maintain high availability during updates, maintenance, or failures
- voluntary disruptions by app owners or cluster admins
 - accidental pod deletion
 - updating pod template causing restart
 - draining a node for repair or upgrade
 - cluster scale down operation

can specify only one of **minAvailable** and **maxUnavailable** in a single PodDisruptionBudget.

```
1 apiVersion: policy/v1
2 kind: PodDisruptionBudget
3 metadata:
4   name: zk-pdb
5 spec:
6   minAvailable: 2
7   selector:
8     matchLabels:
9       app: zookeeper
10
```

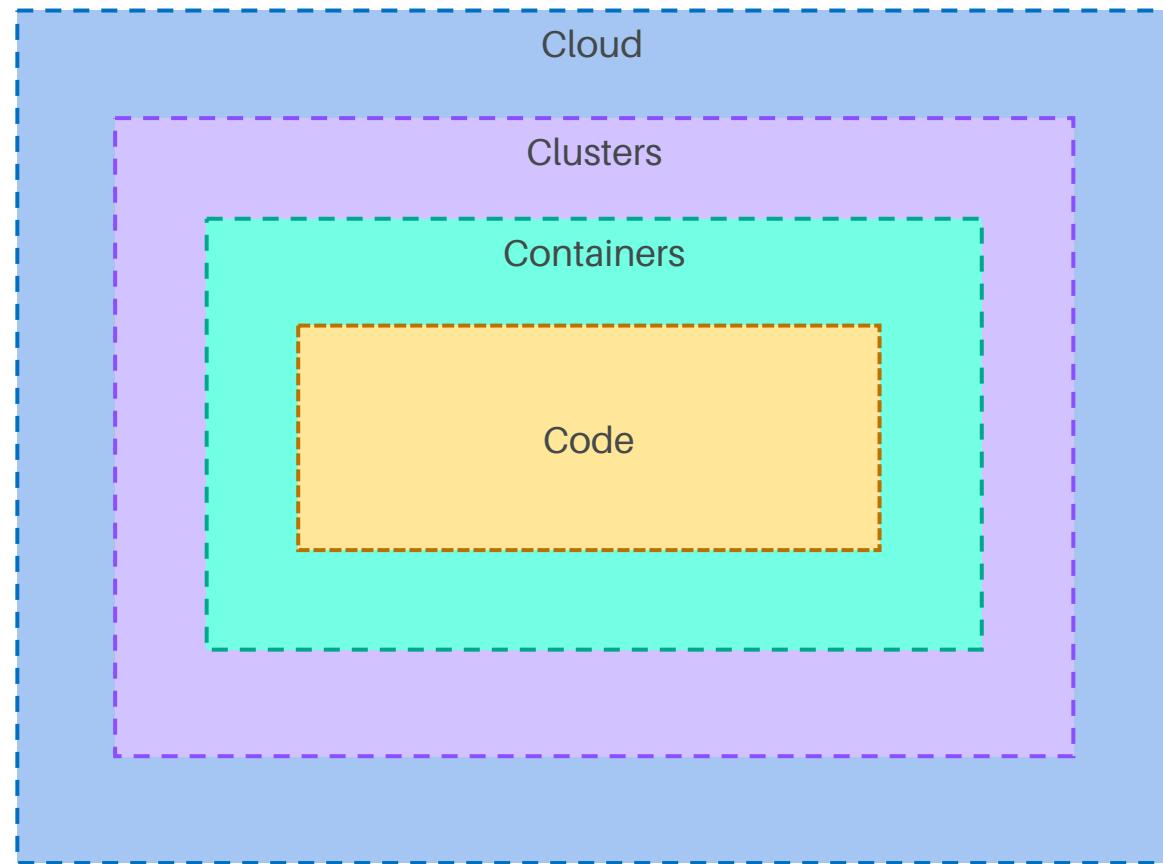
```
1 apiVersion: policy/v1
2 kind: PodDisruptionBudget
3 metadata:
4   name: zk-pdb
5 spec:
6   maxUnavailable: 1
7   selector:
8     matchLabels:
9       app: zookeeper
10
```



Kubernetes Security

RIZMAXed

- Kubernetes is based on cloud-native architecture
- Adheres to CNCF guidelines for best practices in cloud-native information security
- The 4C's of Cloud Native security
 - Cloud
 - Clusters
 - Containers
 - Code
- Falco
 - Cloud native runtime security tool
 - Real-time detection and alerts on abnormal behavior and potential security threats
 - CNCF graduated project





Authentication

RIZMAXed

- Checks who has access to the k8s API
 - process of verifying the identity of entities accessing the cluster
 - users, processes, or components
- Kubernetes doesn't handle user or group management
- Authentication methods
 - X.509 Certificates
 - authenticating users and components using TLS client certificates
 - e.g. kubectl use this method
 - Static tokens
 - using bearer tokens for non-interactive authentication
 - Service Account Tokens
 - for pods and services to access the k8s API
 - allows to authenticate with API server from inside of our pod
 - OIDC Authentication
 - to integrate external identity providers using OpenID Connect
 - uses **JWT** (JSON Web Tokens)

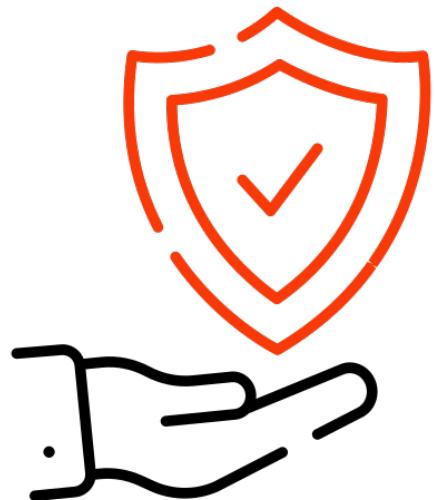




Authorization

RIZMAXed

- Determines what actions the authenticated entity can perform
- Controls access to k8s resources based on defined policies
- Authorization Methods
 - Node
 - controlling access to nodes and their resources
 - ABAC (Attribute Based Access Control)
 - defining access policies based on attributes of users, resources, or environments
 - RBAC (Role Based Access Control)
 - assigning roles to users and service accounts
 - defining permissions based on those roles
 - allows for fine-grained access-control
 - recommended
 - Webhook
 - checking for permissions using a webhook
- OPA Gatekeeper (Open Policy Agent Gatekeeper)
 - a policy engine for Kubernetes, an open-source project
 - admission controller that enforces policies on resources during the admission process





RBAC Authorization

RIZMAXed

- RBAC API declares four kinds of objects
 - Role
 - defines a set of permissions that can be applied to resources within a namespace.
 - ClusterRole
 - defines a set of permissions that apply across the entire cluster.
 - RoleBinding
 - associates a role with a user, group, or service account within a namespace.
 - ClusterRoleBinding
 - associates a cluster-wide role with a user, group, or service account.
- Best practice = rule of **least privilege**
 - grant only the permissions necessary for each user or component to perform their tasks

```
1 apiVersion: rbac.authorization.k8s.io/v1
2 kind: Role
3 metadata:
4   namespace: default
5   name: pod-reader
6 rules:
7   - apiGroups: [""] # "" indicates the core API group
8     resources: ["pods"]
9     verbs: ["get", "watch", "list"]
```

```
1 apiVersion: rbac.authorization.k8s.io/v1
2 kind: RoleBinding
3 metadata:
4   name: read-pods
5   namespace: default
6 subjects:
7   - kind: User
8     name: jane
9     apiGroup: rbac.authorization.k8s.io
10 roleRef:
11   kind: Role
12   name: pod-reader
13   apiGroup: rbac.authorization.k8s.io
```



Hands-on Demo

RIZMAXed

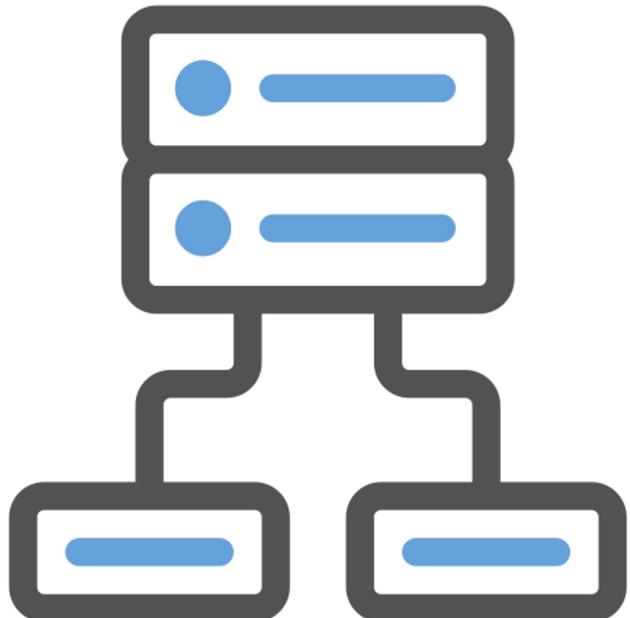




Networking in Kubernetes

RIZMAXed

- Kubernetes uses CNI plugins for cluster networking
 - CNI = Container Network Interface
- Cluster DNS
 - built-in DNS service that provides DNS resolution for services and pods within a cluster
- Each pod gets its own IP address within the cluster.
- Pods can communicate with each other without NAT
- Services
 - allow you to expose an app running in pods
- Ingress
 - used to manage external access to services in a cluster
- Network Policies
 - for controlling traffic between pods and external networks





Network Policies

RIZMAXed

- Rules for controlling traffic flow between pods and external networks
- Specify which pods can communicate with each other and on which ports
- Helps enforce security and segmentation within the cluster.

```
1 apiVersion: networking.k8s.io/v1
2 kind: NetworkPolicy
3 metadata:
4   name: test-network-policy
5   namespace: default
6 spec:
7   podSelector:
8     matchLabels:
9       role: db
10  policyTypes:
11    - Ingress
12    - Egress
13  ingress:
14    - from:
15      - ipBlock:
16        cidr: 172.17.0.0/16
17        except:
18          - 172.17.1.0/24
19      - namespaceSelector:
20        matchLabels:
21          project: myproject
22      - podSelector:
23        matchLabels:
24          role: frontend
25    ports:
26      - protocol: TCP
27        port: 6379
28  egress:
29    - to:
30      - ipBlock:
31        cidr: 10.0.0.0/24
32    ports:
33      - protocol: TCP
34        port: 5978
```



Hands-on Demo

RIZMAXed

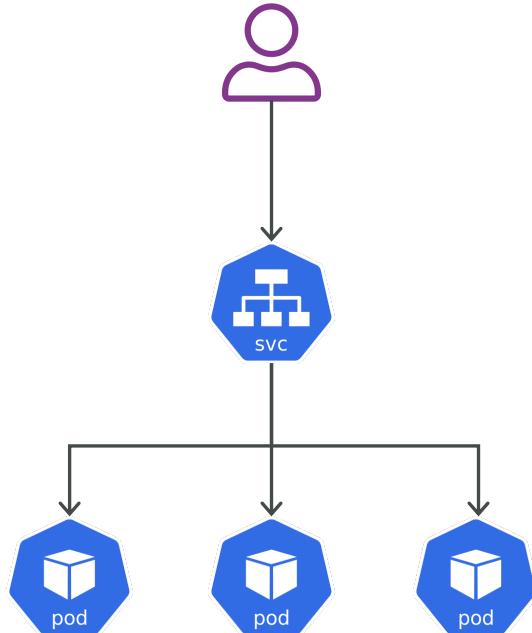




Kubernetes Services

RIZMAXed

- Service allows you to expose an app running in a set of pods
 - provides a stable endpoint to access apps
 - abstracts away the underlying pod IP addresses
 - uses labels and selectors to route traffic to pods
 - supports load balancing and service discovery
- Service types
 - Cluster IP (default)
 - exposes service internally on an IP within cluster
 - NodePort
 - exposes service externally on a static port on each node
 - LoadBalancer
 - exposes service externally using an external load balancer
 - e.g. AWS ELB
 - ExternalName
 - exposes service externally using provided external name or hostname
 - maps a service to a DNS hostname using a CNAME record
- Headless Service (clusterIP=None)
 - Cluster IP service without an IP address
 - used with StatefulSets



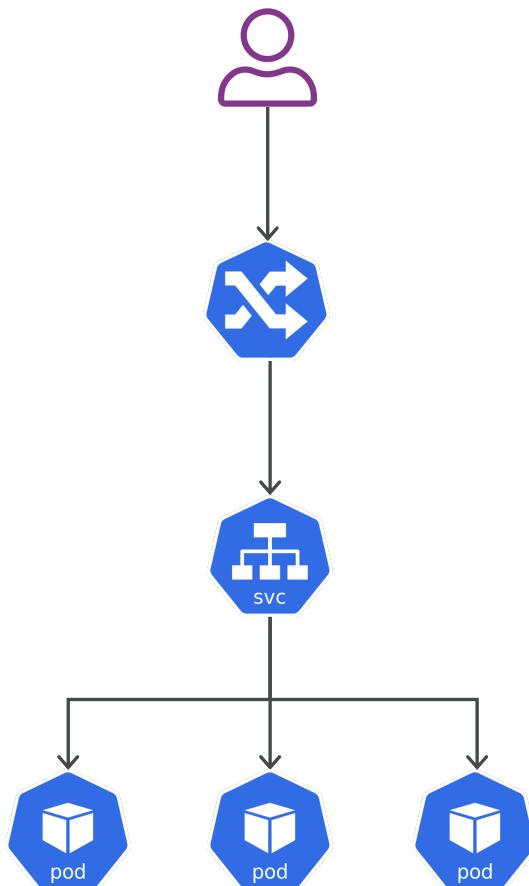
If you create a service without selectors, you must create corresponding endpoint objects manually.



Ingress

RIZMAXed

- Used to manage external access to services in a cluster
- Specifically, HTTP/S applications, websites and APIs
- Supports
 - TLS termination
 - path-based routing
 - virtual hosting



```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: example-ingress
spec:
  tls:
    - hosts:
        - example.com
      secretName: example-tls
  rules:
    - host: example.com
      http:
        paths:
          - path: /app
            pathType: Prefix
            backend:
              service:
                name: webapp-service
                port:
                  number: 80
          - path: /api
            pathType: Prefix
            backend:
              service:
                name: api-service
                port:
                  number: 80
```



Hands-on Demo

RIZMAXed



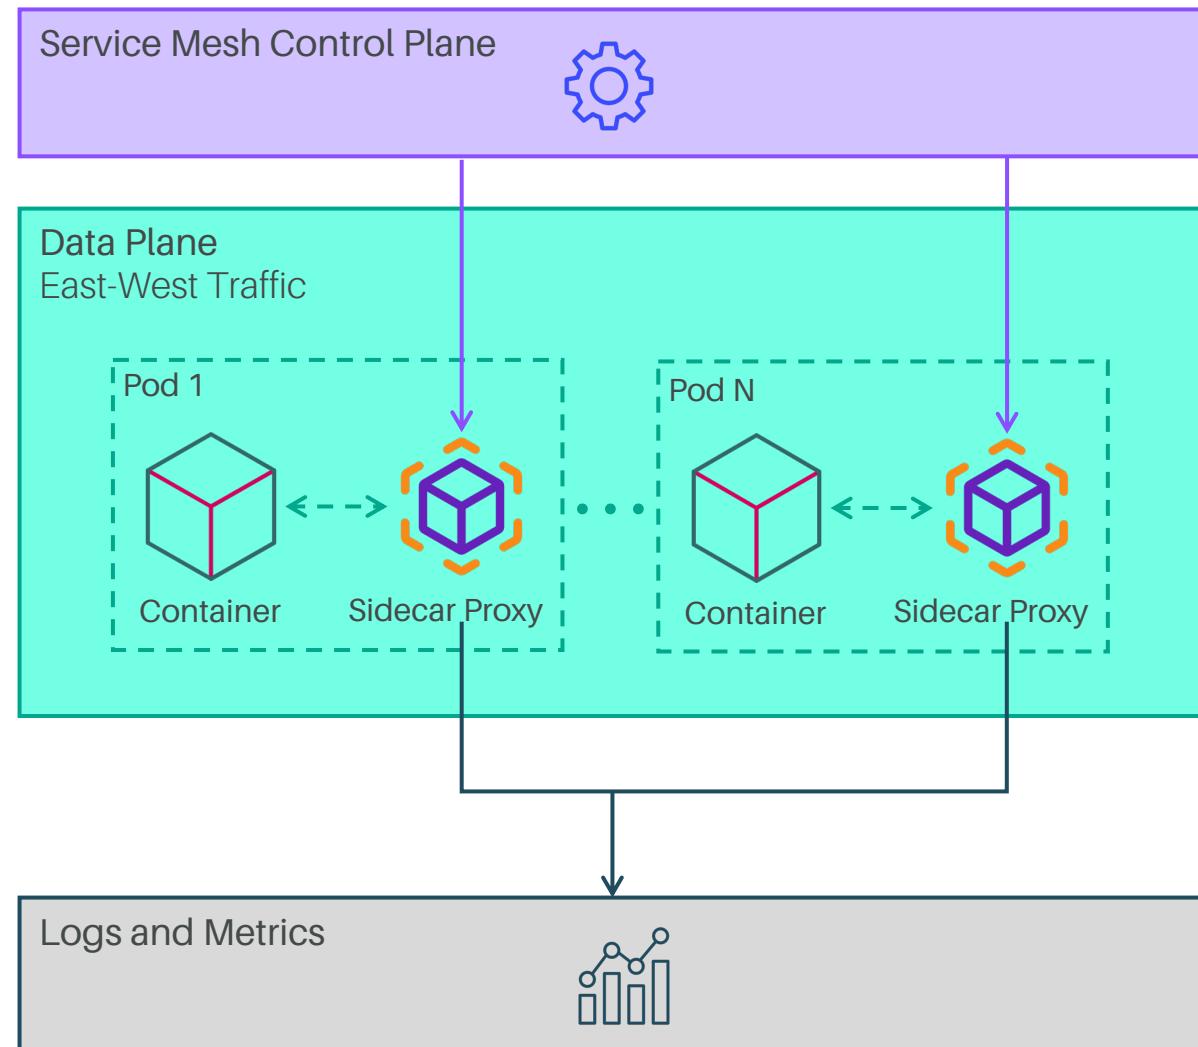


Service Mesh

RIZMAXed

- Dedicated infrastructure layer
- Manages communication between microservices within a cluster
- Provides traffic management, observability, and security features
 - routing, load balancing, and retries
 - centralized logging, and metrics
 - implements security policies for inter service communication
- Uses sidecar proxies to intercept and control network traffic
- Kubernetes provides SMI
 - **SMI = Service Mesh Interface** 
- Popular solutions that support SMI
 - **Linkerd, Istio, Consul** 

Service Mesh Architecture





Kubernetes Storage

RIZMAXed

- Ephemeral volume types have a lifetime of a pod
- Persistent volumes exist beyond the lifetime of a pod
- emptyDir volume
 - initially empty, used as a scratch disk
 - all containers in the pod can read/write the same files in the volume
 - erased when a Pod is removed
- iscsi volume
 - to mount an existing iscsi volume
 - contents are preserved when pod is removed
- nfs volume
 - to mount an existing NFS share
 - contents are preserved when pod is removed

```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: test-pd
5  spec:
6    containers:
7      - image: registry.k8s.io/test-webserver
8        name: test-container
9        volumeMounts:
10       - mountPath: /cache
11         name: cache-volume
12    volumes:
13      - name: cache-volume
14      emptyDir:
15        sizeLimit: 500Mi
```



Persistent Volumes

RIZMAXed

- PV = PersistentVolume
 - k8s objects representing storage volumes within the cluster
- PVC = PersistentVolumeClaim
 - k8s object used to request storage from a PV
 - binds a PV and allows you to mount it on a pod
- Reclaim Policy defines action to take on PV when PVC is deleted
 - Retain - manual reclamation
 - Recycle - automatic reclamation using basic scrub (`rm -rf /thevolume/*`)
 - Delete - delete the volume
- Storage Classes
 - define the type and properties of storage volumes available within the cluster





Persistent Volumes - Example

RIZMAXed

```
1 apiVersion: v1
2 kind: PersistentVolume
3 metadata:
4   name: example-pv
5 spec:
6   capacity:
7     storage: 1Gi
8   volumeMode: Filesystem
9   accessModes:
10    - ReadWriteOnce
11 persistentVolumeReclaimPolicy: Retain
12 storageClassName: manual
13 hostPath:
14   path: /data/example-pv
15
```



```
1 apiVersion: v1
2 kind: PersistentVolumeClaim
3 metadata:
4   name: example-pvc
5 spec:
6   accessModes:
7     - ReadWriteOnce
8   resources:
9     requests:
10      storage: 1Gi
11   storageClassName: manual
12
```



```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: example-pod
5 spec:
6   containers:
7     - name: app-container
8       image: nginx
9       volumeMounts:
10        - name: data-volume
11          mountPath: /data
12   volumes:
13     - name: data-volume
14       persistentVolumeClaim:
15         claimName: example-pvc
16
```





Container Storage Interface

RIZMAXed

- Kubernetes provides CSI to interact with storage systems
 - **CSI = Container Storage Interface**
- allows integration between containerized apps and storage platforms
- Popular solutions that support CSI 
 - **Rook**
 - CNCF graduated project
 - Storage Orchestration for Kubernetes with support for **Ceph** storage
 - automates deployment and management of Ceph
 - provide self-managing, self-scaling, and self-healing storage
 - **Longhorn**
 - CNCF graduated project
 - Cloud-native distributed storage for Kubernetes



Autoscaling K8S Workloads

RIZMAXed

- HPA = Horizontal Pod Autoscaler
 - scales pod replicas automatically
 - based on CPU utilization or other metrics
- VPA = Vertical Pod Autoscaler
 - scales CPU and memory reservations of your pods automatically
 - based on resource utilization
- Cluster Autoscaler
 - scales # of nodes in the cluster automatically
 - based on resource demands

```
apiVersion: autoscaling.k8s.io/v1
kind: VerticalPodAutoscaler
metadata:
  name: myapp-vpa
spec:
  targetRef:
    apiVersion: "apps/v1"
    kind: "Deployment"
    name: "myapp-deployment"
  updatePolicy:
    updateMode: "Auto"
```

```
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: php-apache
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: php-apache
  minReplicas: 1
  maxReplicas: 10
  metrics:
    - type: Resource
      resource:
        name: cpu
      target:
        type: Utilization
        averageUtilization: 50
```



Autoscaling K8S Workloads (contd.)

RIZMAXed

- KEDA
 - Kubernetes Event-driven Autoscaling
 - KEDA is a CNCF graduated project
 - scale automatically based on event sources
 - e.g. scale based on # of messages in a queue
 - Azure Queue, Kafka topics, or AWS SQS
- Karpenter
 - open-source cluster Autoscaler
 - created by AWS
 - automatically provisions new nodes in response to unschedulable pods
- Autoscaling Best Practices
 - Use the Right Metrics
 - Test Your Autoscaling
 - Monitor and Adjust

```
1 apiVersion: keda.sh/v1alpha1
2 kind: ScaledObject
3 metadata:
4   name: sqs-scaledobject
5 spec:
6   scaleTargetRef:
7     deploymentName: myapp-deployment
8     pollingInterval: 5
9     minReplicaCount: 1
10    maxReplicaCount: 10
11    triggers:
12      - type: aws-sqs
13        metadata:
14          queueName: my-sqs-queue
15          awsRegion: us-west-2
16          identityResourceID: my-aws-iam-role-arn
17          queueLength: "5"
```



Hands-on Demo

RIZMAXed





Helm and Helm Charts

RIZMAXed

- Helm
 - package manager for Kubernetes
 - CNCF graduated project
 - simplifies the process of deploying, managing, and upgrading applications
- Helm Chart
 - packaged Kubernetes resources
 - reusable, versioned bundles containing
 - templates for Kubernetes manifests
 - default config
 - optional dependencies
- How it works
 - Developers create Helm charts to package k8s app
 - Users install Helm chart onto their k8s cluster using `helm install`
 - User can manage k8s apps using `helm upgrade`, and `helm rollback`



```
$ helm install my-web-app ./path/to/chart
```

```
$ helm upgrade my-web-app ./path/to/updated/chart
```

```
$ helm rollback my-web-app
```



Cloud Native Computing

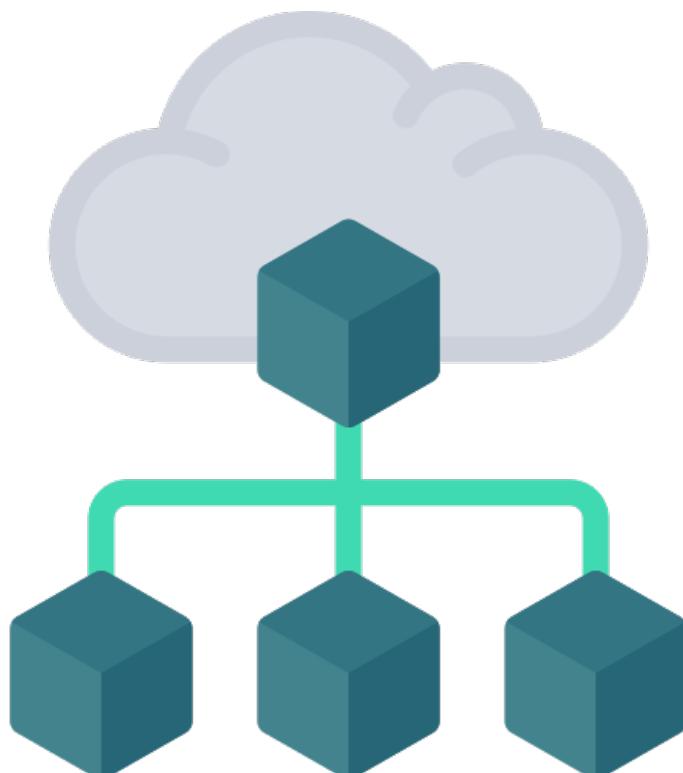
Cloud Native, Serverless, and Microservices



Cloud Native Computing

RIZMAXed

- = design and deployment of applications that fully leverage the benefits of cloud computing models (public, private, or hybrid)
- => scalable, resilient, agile, cost-effective, portable applications
- Key guiding principles
 - containerization and orchestration
 - microservices architecture
 - continuous integration/delivery/deployment
 - DevOps, GitOps, FinOps
- Use cases
 - distributed web apps
 - IoT and edge computing
 - big data and analytics

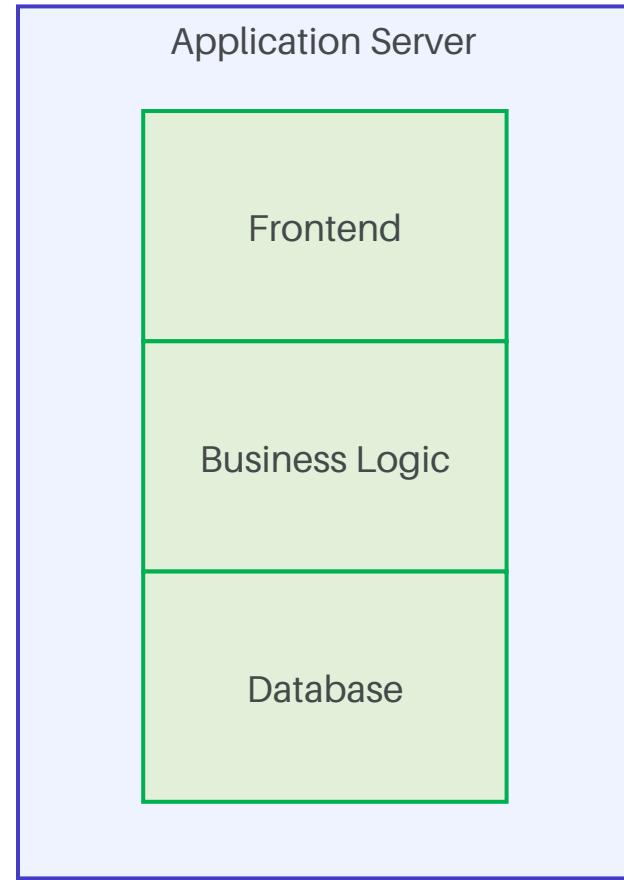
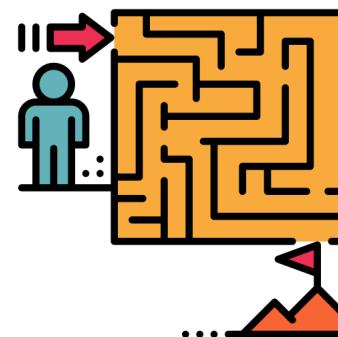




Traditional Approach to Software Development

RIZMAXed

- Monolithic architecture
 - All components run as a closely-knit unit => tightly-coupled
 - Not highly scalable
 - Not flexible
 - Not agile
 - Complex deployment (not automated)
 - Difficult to maintain

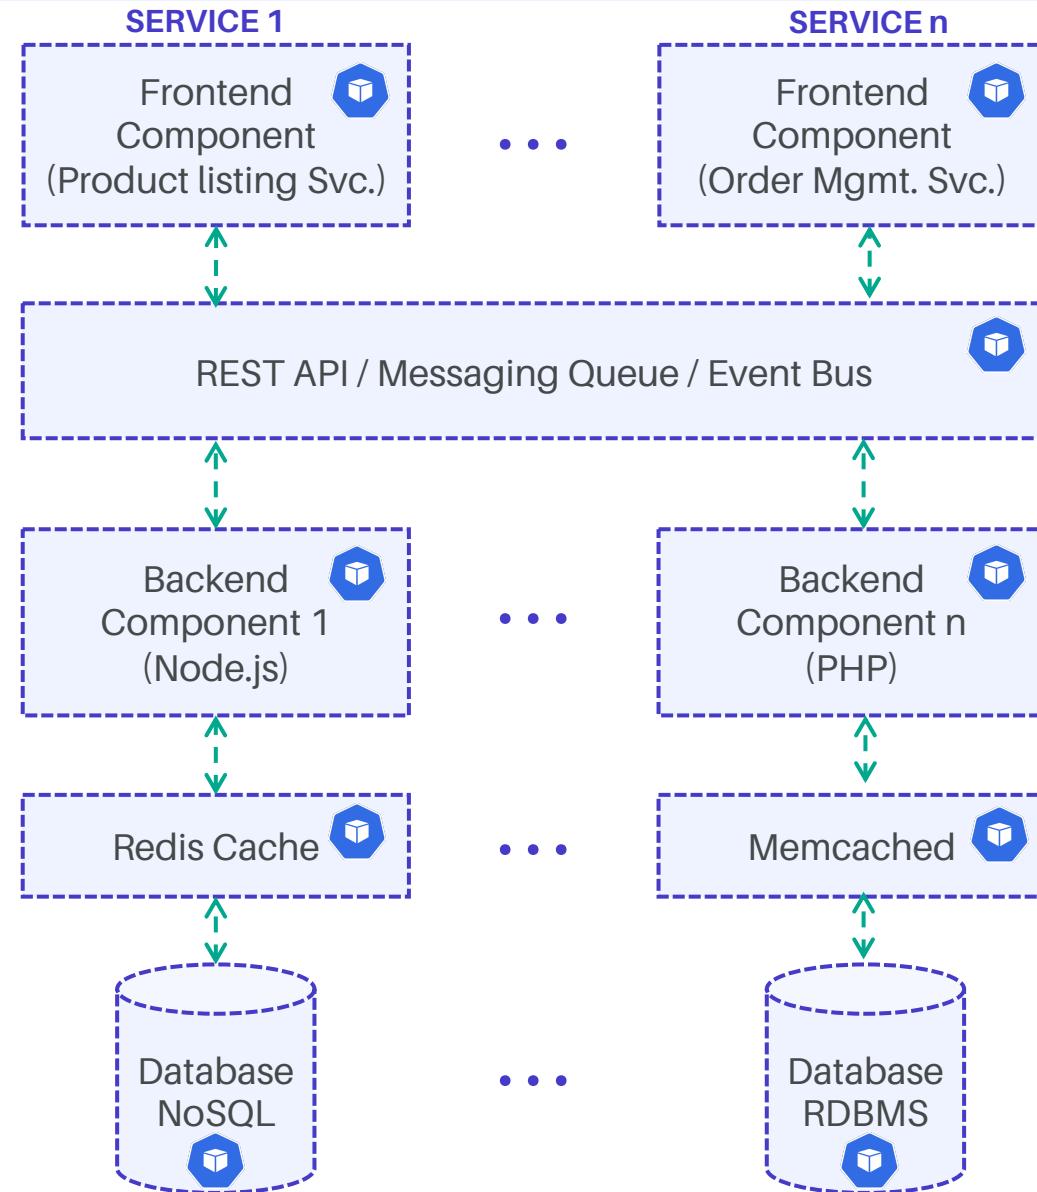




Microservices Architecture

RIZMAXed

- Structures an application as a collection of services
- Each service is
 - independent unit
 - can be deployed on its own
 - often owned by a single, small team
 - highly and independently scalable
 - supports automated deployments
 - independently maintainable
- Loosely coupled architecture
 - allows for distributed computing
- Flexible, agile (allows for rapid, frequent changes)

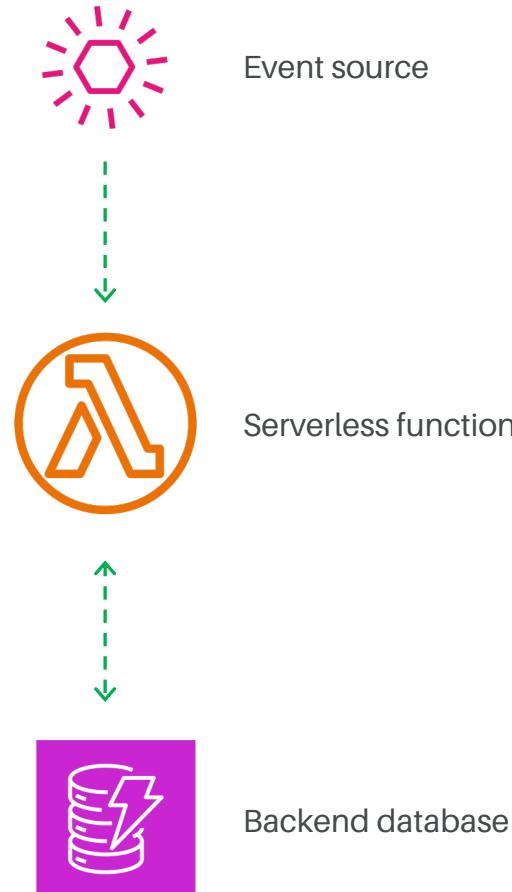




Serverless Computing

RIZMAXed

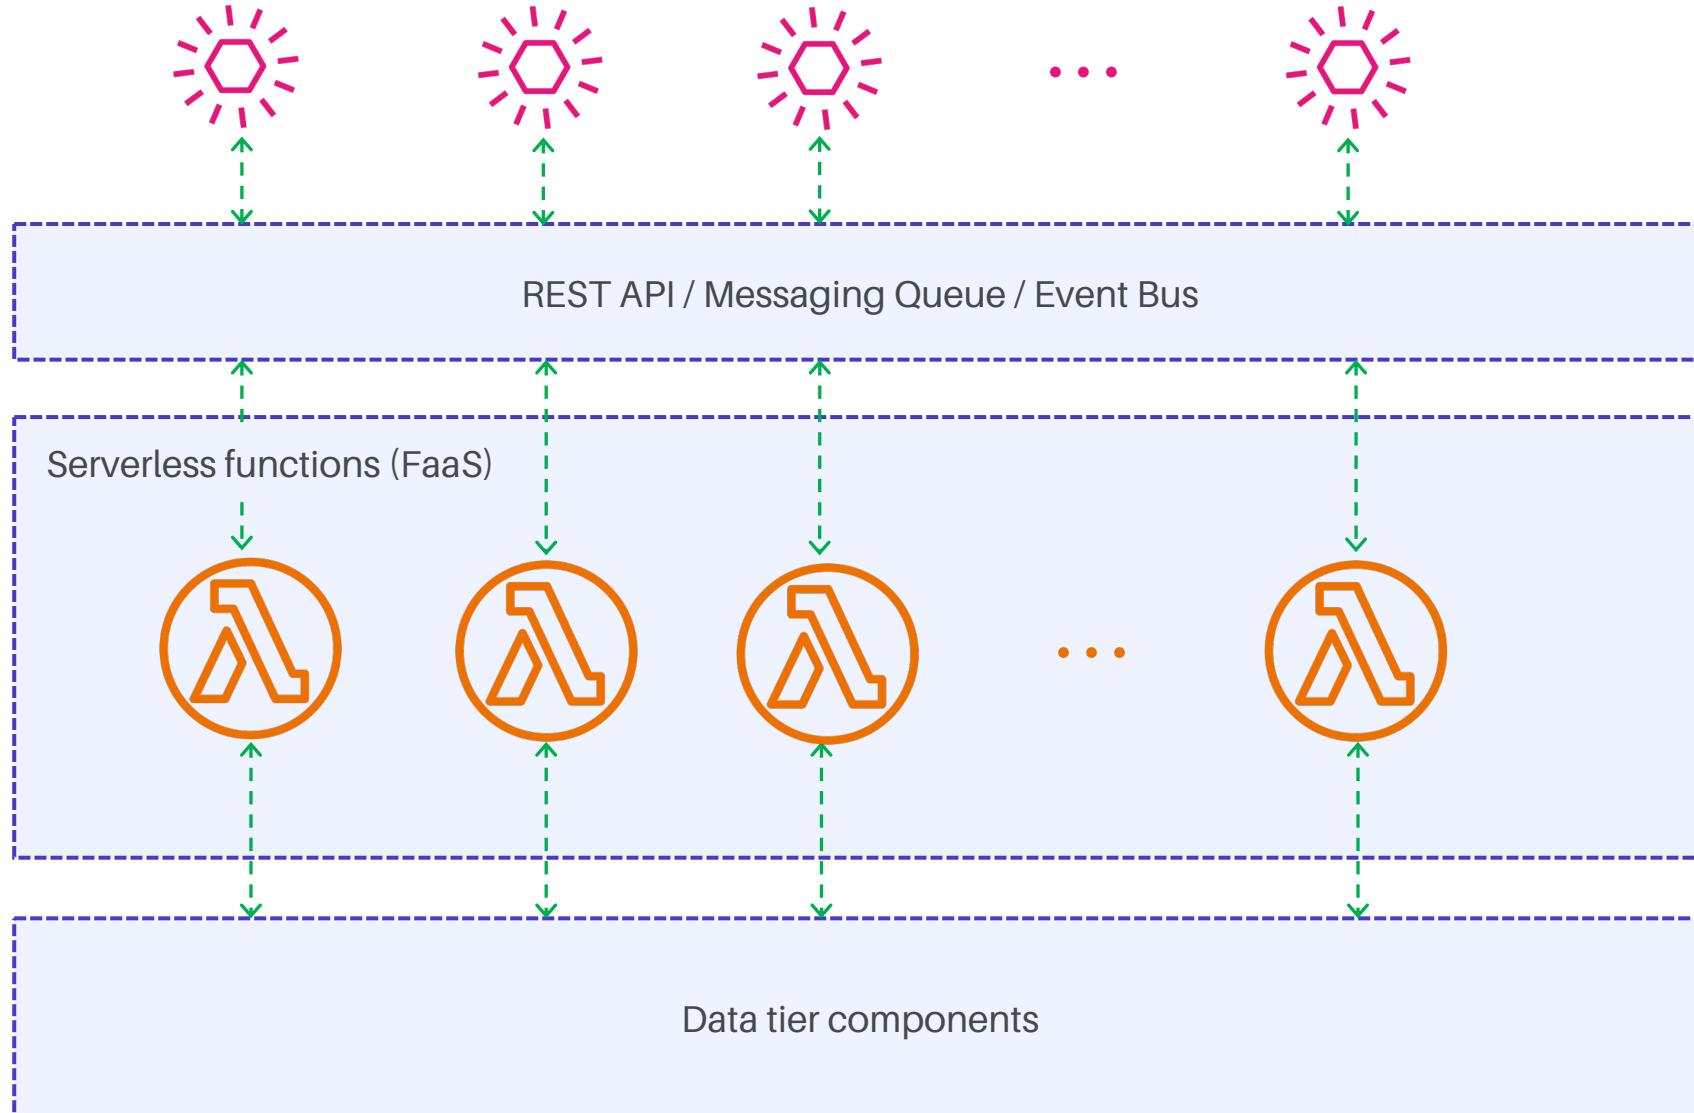
- Cloud provider manages servers
- Developers focus on writing code (functions)
 - without worrying about underlying infrastructure
 - => FaaS (Function as a Service)
- Serverless functions
 - event-driven
 - auto-scaling
 - pay-per-use
- Benefits
 - No servers to manage
 - Cost-effective
 - High scalability
 - Fast time-to-market
- FaaS Solutions on public cloud
 - AWS Lambda
 - Azure Functions
 - Google Cloud Functions





Serverless Microservices Architecture

RIZMAXed





Kubernetes Serverless Solutions

RIZMAXed

- Knative
 - serverless application layer built on top of Kubernetes
 - three components
 - Knative Serving - HTTP-triggered autoscaling container runtime
 - Knative Eventing - CloudEvents-over-HTTP asynchronous routing layer
 - Knative Functions - Serverless function framework
- OpenFaaS
 - serverless framework for Kubernetes and Docker Swarm
 - platform for building serverless functions with containers
 - write services and functions in multiple languages
 - build and ship your code in a container image





Community and Governance

RIZMAXed

- CNCF = Cloud Native Computing Foundation
 - is a vendor-neutral organization
 - that focuses on growth and adoption of cloud-native technologies
 - through open-source collaboration and community engagement
- Community
 - diverse ecosystem of contributors, users, and enthusiasts worldwide
 - collaborative development, knowledge sharing, and innovation
 - regular events, meetups, and conferences for networking, learning and engagement
 - e.g. KubeCon + CloudNativeCon, Contributor Summits
- Governance
 - open feedback channels enable continuous improvement
 - transparent and inclusive governance models
 - decisions are made openly and democratically
 - enforced code of conduct





CNCF Committees and Groups

RIZMAXed

- GB = Governing Board
 - responsible for marketing, budget and other business oversight decisions
- TOC = Technical Oversight Committee
 - defines and maintains the technical vision
 - provides technical leadership and oversight for CNCF and its projects
 - reviews/approves new projects
 - governance discussions and decision-making
- SIG = Special Interest Group
 - focus on specific domains, technologies, or use cases within cloud native ecosystem
- WG = Working Group
 - SIGs may organize working groups to address specific challenges/initiatives
- EUC = End User Community
 - End User SIGs and User Groups

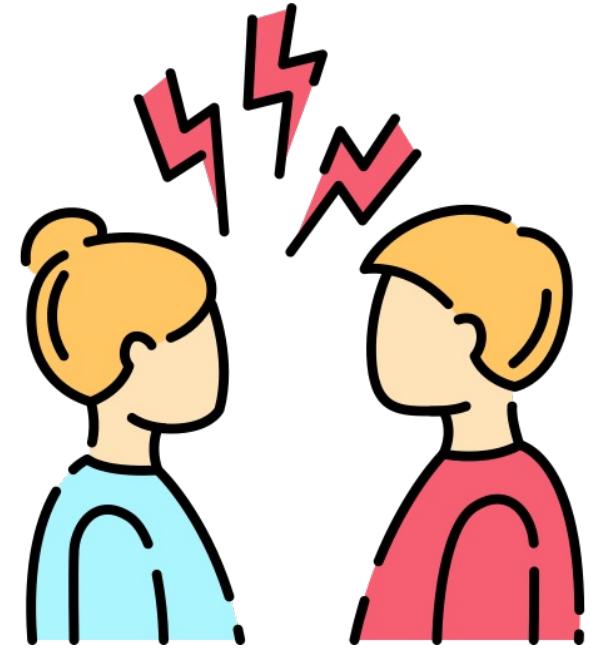




Conflict resolution within CNCF

RIZMAXed

- Open Discussion 
 - open and transparent communication within its groups
- Mediation
 - group facilitators/leaders may act as mediators to discuss and reach a resolution
- Escalation
 - unresolved conflicts may be escalated to TOC, GB for review and resolution
- Consensus building
 - among group members and stakeholders
 - through **elections and voting** 
- Code of Conduct enforcement
 - to ensure a respectful and inclusive environment

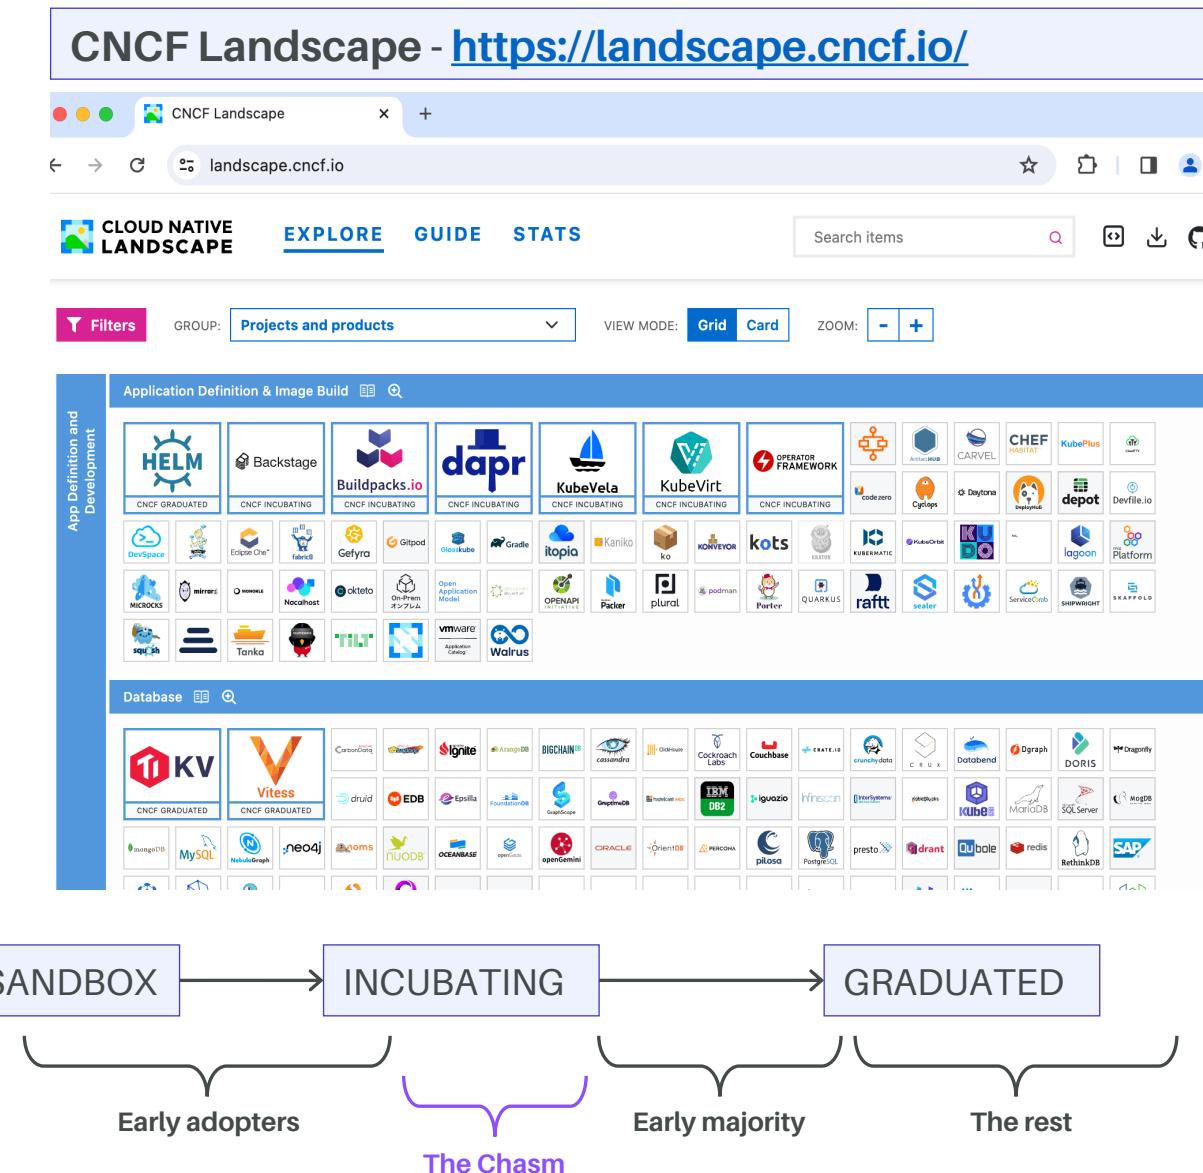




CNCF Landscape

RIZMAXed

- Comprehensive visualization of the cloud-native ecosystem
- Showcases projects and tools
 - exploration and discovery
 - reference
 - community engagement
- Well-defined categories and filtering options
- CNCF projects are tagged as per their maturity levels
 - Sandbox
 - Incubating
 - Graduated
- Crossing the Chasm
 - => moving from early adopters to mainstream adoption
 - i.e. from Incubating to Graduated





Open Standards

RIZMAXed

- OCI = Open Container Initiative
 - Defines open standards around container formats and runtimes
 - Runtime Spec (based on runc)
 - Image Spec
 - Distribution Spec
- CRI = Container Runtime Interface
- CNI = Container Network Interface
- CSI = Container Storage Interface
- SMI = Service Mesh Interface
- Enable diverse stakeholders to collaborate and innovate
- Interoperability
 - between different systems and components within the cloud-native ecosystem
- Portability
 - across different environments and platforms
 - reduces vendor lock-in
 - increases flexibility





Cloud Native Roles and Persons

RIZMAXed

- Application Developer
- DevOps Engineer
- Site Reliability Engineer (SRE)
- Security Engineer
- DevSecOps Engineer
- Cloud Architect
- FinOps Consultant





Application Developer

RIZMAXed

- Responsibilities
 - Design, develop, deploy cloud-native apps
 - using modern technologies like microservices, containers, and serverless
 - using agile practices
- Skills
 - Programming languages
 - Containerization
 - Git and version-control systems
 - CI/CD pipelines





DevOps Engineer

RIZMAXed

- Responsibilities
 - Automate the deployment, scaling, and management
 - of cloud-native apps and infrastructure
 - using infrastructure-as-code (IaC)
 - using CI/CD pipelines
 - using GitOps tools
- Skills
 - Cloud platforms
 - Container orchestration
 - Configuration management
 - Scripting languages





Site Reliability Engineer (SRE)

RIZMAXed

- Crossover role with DevOps
 - with focus on reliability and operational excellence
- Responsibilities
 - Ensure reliability, availability, and performance
 - of cloud-native applications and infrastructure
 - through monitoring, alerting, incident response, and capacity planning
 - Meet SLAs, with SLOs, SLIs
 - SLA = Service Level Agreement
 - LO = Service Level Objective
 - SLI = Service Level Indicator
- Skills
 - Observability and Telemetry tools
 - Incident management
 - Performance optimization
 - Distributed systems architecture





SLA, SLO, and SLI

RIZMAXed

- SLA = Service Level Agreement
 - contract with the client defining the level of service
 - based on measurable metrics
 - uptime, availability, performance, and response times
 - e.g. service uptime SLA = 99.95%
- SLO = Service Level Objective
 - defines the acceptable level of service
 - goals defined to meet each measurable metric in the SL
 - e.g. SLO = 99.97% uptime
- SLI = Service Level Indicator
 - Measure of actual performance against the SLOs
 - Used to monitor and track SLA compliance
 - e.g. SLI = 99.99% uptime



SLA compliance = $SLI \geq SLO$



Security Engineer

RIZMAXed

- Responsibilities
 - Secure cloud-native applications and infrastructure
 - Implement security best practices
 - Implement access control mechanisms
 - Threat detection and vulnerability management
- Skills
 - Cloud native security principles
 - Encryption
 - Identity and access management (IAM)
 - Security compliance frameworks
 - Penetration testing





DevSecOps Engineer

RIZMAXed

- Responsibilities

- Integrate security practices into the DevOps workflow
- Implement security controls, tools, and processes
- Collaborate with dev, ops, and security teams
- Automate security testing, compliance checks, and remediation tasks
- Monitor and analyze security metrics, logs, and alerts

- Skills

- Security principles, standards, and best practices
- Experience with DevOps practices
- Scripting / programming skills
- Cloud platforms





Cloud Architect

RIZMAXed

- Responsibilities
 - Design and implement cloud-native architectures
 - to meet scalability, reliability, and performance requirements
 - using cloud services, microservices, hybrid cloud solutions
- Skills
 - Cloud platforms
 - Networking
 - Containerization
 - Distributed systems
 - Architecture patterns
 - System Design skills
 - Communication skills





FinOps Consultant

RIZMAXed

- Responsibilities
 - Cloud cost optimization
 - Analyze cloud spending trends, identify cost-saving opportunities
 - Recommend strategies to improve cost-efficiency
 - Diagnose cost inefficiencies, identify root causes
 - Develop practical solutions to optimize cloud spending
- Skills
 - Financial management skills
 - budgeting, forecasting, and cost analysis
 - In-depth knowledge of cloud platforms
 - Familiarity with cloud billing and pricing mechanisms
 - Cloud cost management tools





Telemetry & Observability

Prometheus, Grafana, and FinOps



Telemetry and Observability

RIZMAXed

- Telemetry
 - collect, aggregate, and transmit data outputs from a system
 - for monitoring and performance analysis
- Observability
 - infer the internal state and behavior of a system
 - based on its external outputs i.e. using telemetry data
- Key components
 - Metrics
 - measurements over a time period e.g. CPU, memory utilization etc.
 - Categories – **Meters, Counters, Gauges, Histograms**
 - Logs
 - recorded messages generated by the systems and components
 - based on verbosity like error, warning, info, debug etc
 - Trace
 - represents journey of a request as it moves through a distributed system
 - made up of a series of tagged time intervals called **spans**
 - Events
 - Notifications and alerts generated by a system

OBSERVABILITY TOOLS

Prometheus

Monitoring & alerting toolkit

Grafana

Visualization and dashboarding platform

ELK stack

Elasticsearch, Logstash, Kibana
Stack for collecting, indexing, and visualizing logs

Jaeger

Distributed tracing solution

Zipkin

Distributed tracing solution



- Open-source monitoring and alerting toolkit, CNCF graduated
- Widely used for monitoring Kubernetes
- Originally built at SoundCloud
- Collects and stores **time-series data** in a database
 - Can be queried using **PromQL** (Prometheus Query Language)
- **Prometheus Server**
 - Discovers targets (pods, services) using Kubernetes service discovery
 - Collects data from targets using **HTTP scraping (pull-based)**
- **Metrics Exporter**
 - gathers metrics from systems that do not natively expose Prometheus-compatible endpoints
 - agents deployed alongside Kubernetes applications
 - to expose metrics in Prometheus-compatible format (**push based**)
- **Alertmanager**
 - Handles alerts generated by Prometheus
 - Sends notifications via various channels (email, Slack, PagerDuty)



- Open-source analytics and visualization platform
- Widely used for visualizing metrics and logs in real-time
 - System performance and health
- Flexible, interactive, customizable dashboards and panels
- Supports a wide range of data sources
 - Prometheus, Elasticsearch, InfluxDB, etc.
- Can set up alerts based on predefined thresholds and conditions
- Supports templating for dynamic filtering and grouping of data



Cost Management in Kubernetes

RIZMAXed

- Refers to optimizing costs on cloud resources
- Right sizing – use right sized resources e.g. CPU, memory
- Auto scaling – dynamically adjust resource capacity
- Resource tagging – to track and allocate costs accurately
- Cost allocation
 - allocating costs to teams, projects, or departments
 - for accountability and budgeting purposes
- Reserved instances
 - securing discounted pricing using long-term reservations
- Spot instances or preemptible VMs
 - to take advantage of discounted pricing
 - for non-critical workloads or batch processing tasks
- Cloud Anomaly Detection
 - to identify abnormal patterns or deviations in cloud infrastructure, apps and services
 - anomalies in resource utilization, or billing discrepancies

COST MANAGEMENT TOOLS

AWS Cost Explorer

AWS tool for visualizing and analyzing AWS cost and usage

Azure Cost Management

Azure's cost monitoring and optimization service

Google Cloud Billing

collection of tools to track and your Google Cloud spending

Kubecost

Cost monitoring for Kubernetes workloads and cloud costs

CloudHealth

Multi-cloud FinOps platform



Cloud Native Application Delivery

CICD, GitOps, and more...



Cloud Native Application Delivery

RIZMAXed

- Modern approach to building, deploying, and managing applications
- What?
 - Containerization and orchestration
 - Microservices architecture
 - Infrastructure as Code (IaC)
 - CI/CD automation and GitOps





Cloud Native Application Delivery

RIZMAXed

- Why?
 - Agility
 - rapid iteration, experimentation, and innovation
 - Scalability
 - apps can handle varying workloads and demands
 - Resilience
 - reliability with fault tolerance, and minimal disruptions
 - Cost-efficiency
 - pay-as-you-go, pay for what you use





Cloud Native Application Delivery

RIZMAXed

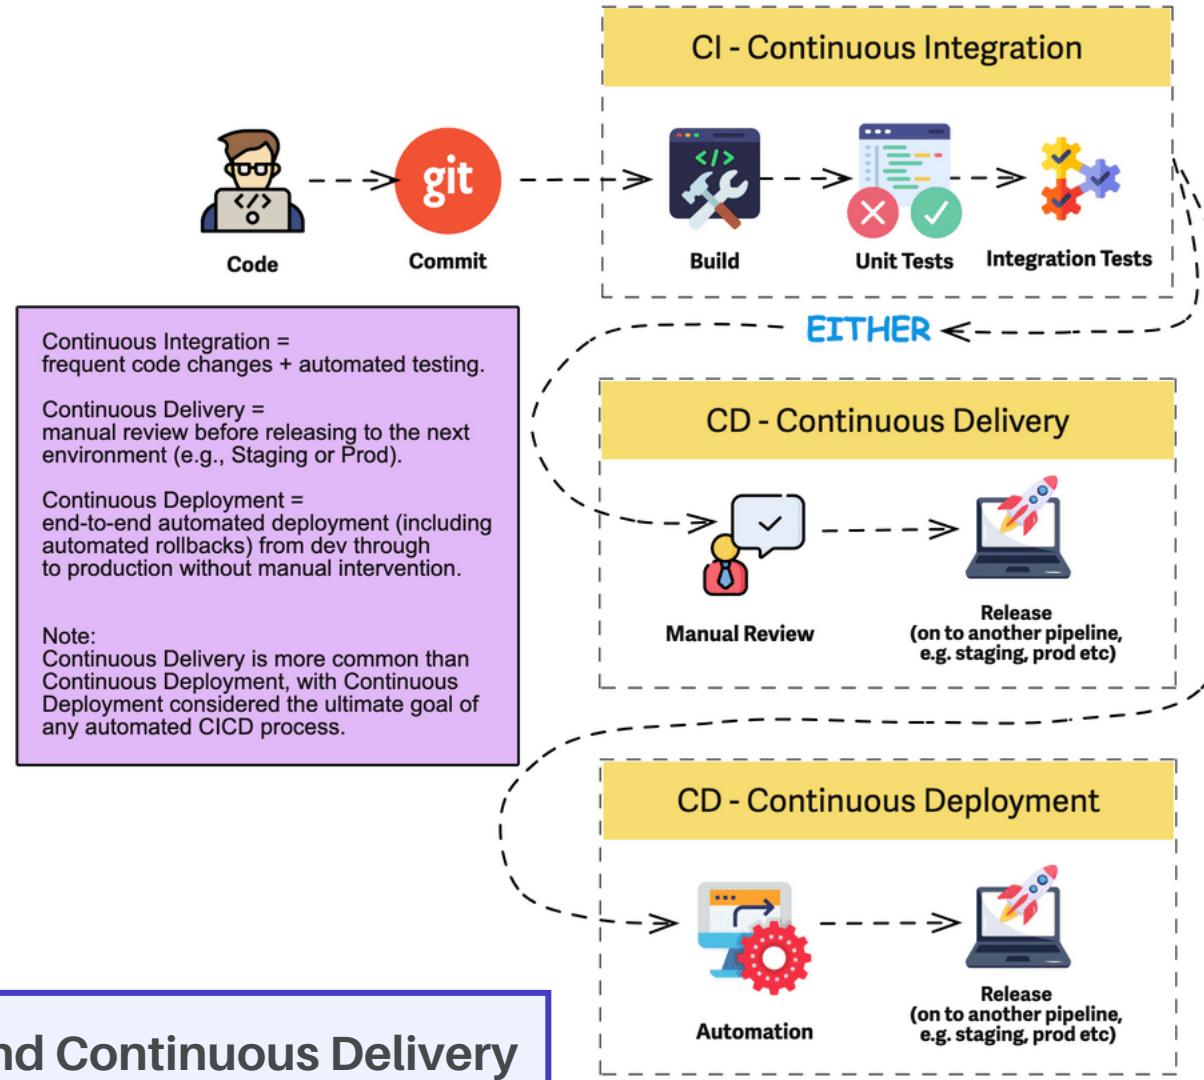
- How?
 - Embrace agile methodologies
 - Design for resilience and scale
 - Embrace DevOps culture
 - Automate everything
 - Monitor and iterate
 - Telemetry and Observability



Continuous Integration -> Continuous Delivery -> Continuous Deployment



- **Automate build, test, and deploy process**
- Continuous Integration (CI)
 - automatically **build and test** code changes on commit to Git repo
- Continuous Delivery (CD)
 - automatically **deploy to staging or pre-production**
- Continuous Deployment (CD)
 - automatically **deploy to production** (without manual intervention)



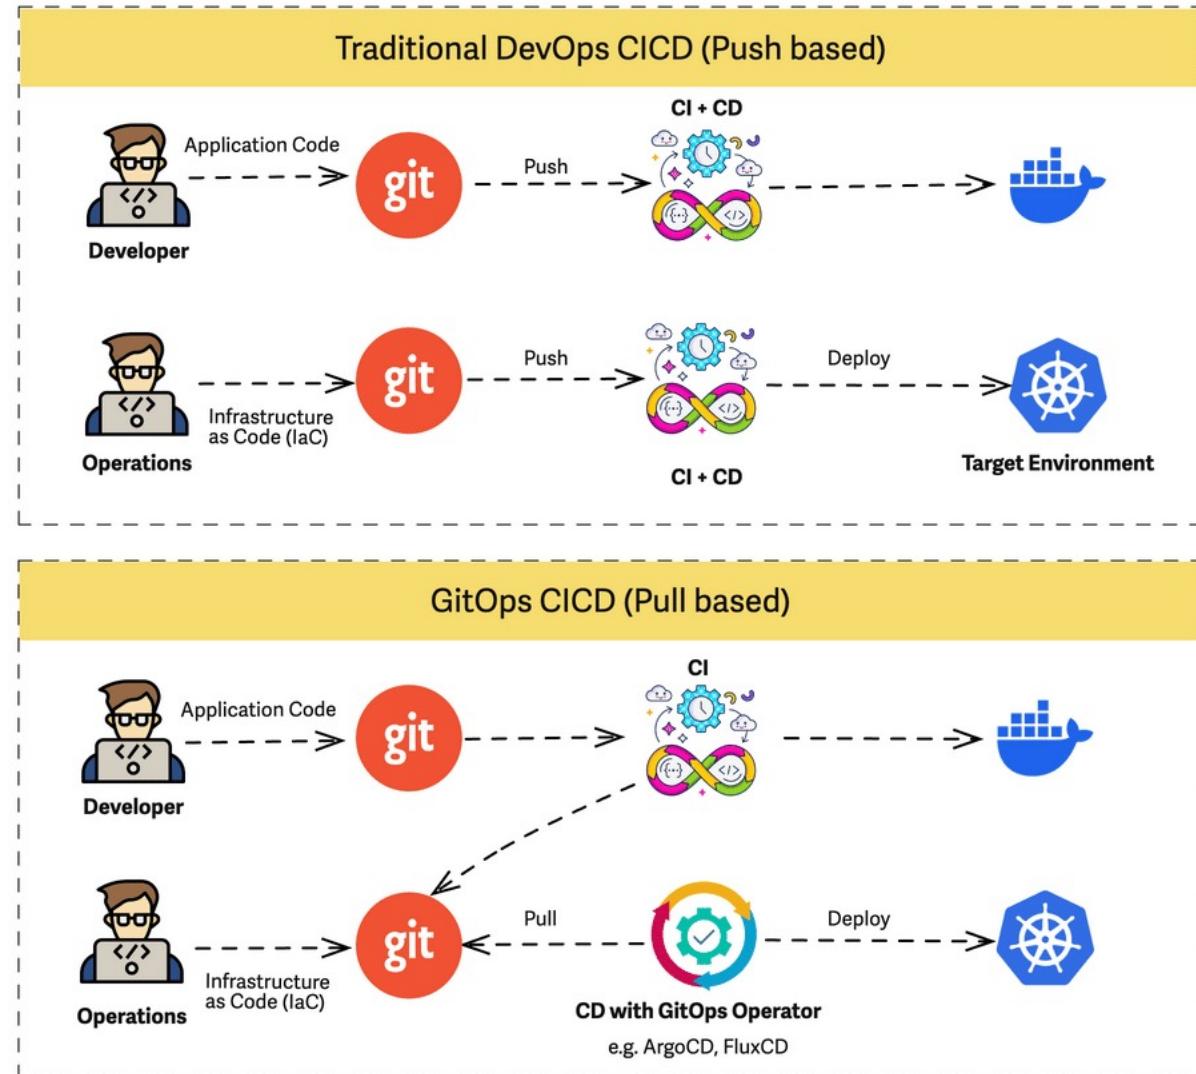
CI/CD typically refers to: **Continuous Integration and Continuous Delivery**



GitOps

RIZMAXed

- DevOps
 - development (Dev) and operations (Ops) teams work together to deliver apps using deployment automation
 - CI/CD automation, IaC, Observability
 - Traditionally push-based deployment
- GitOps
 - extends DevOps principles
 - focus on automated continuous sync
 - uses a Git operator to reconcile
 - desired state in Git
 - with actual state of the cluster
 - GitOps operators – ArgoCD, FluxCD





ArgoCD and FluxCD

RIZMAXed

- ArgoCD and FluxCD are two popular GitOps Operators
- ArgoCD enables GitOps workflows
 - continuously monitors Git repositories for changes
 - automatically applies them to Kubernetes clusters
 - => ensures desired state convergence
 - provides Image Updater
 - to automatically update the container images used in k8s manifests
- FluxCD is built with the GitOps Toolkit
 - GitOps Toolkit is collection of open-source tools and best practices to implement GitOps workflows
 - continuously monitors git and applies changes to k8s
 - Image reflector and automation controllers
 - work together to update Git when container images are updated





Tackling the exam questions

"Perfect" practice makes perfect!



Exam Tips

RIZMAXed

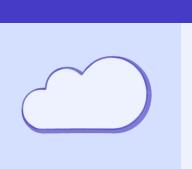
- Focus on the question first, before reading the responses
 - try to answer the question if possible
 - then see if there is a matching response
- Eliminate choices containing obviously incorrect statements, and using your knowledge
- Identify key phrases + look for qualifiers in the questions
 - the most operationally efficient
 - the most cost-effective option
 - the fastest
 - the least administrative overhead etc.



Exam Tips

RIZMAXed

- If you can't figure out the correct answer
 - make a best guess (there is no penalty for incorrect answers, answer all the questions)
 - flag the question for later review
 - move on to the next question



Course Roundup

Woohoo! You did it!



Congratulations



RIZMAXed

- You're all set!
- Revisit the lectures, review the course slides
- practice hands-on demos if you like
- Go ahead and take the final exam!



Some more resources

RIZMAXed

- Review Kubernetes documentation wherever possible
 - <https://kubernetes.io/docs/home/>
- Review the kubectl Cheat Sheet
 - <https://kubernetes.io/docs/reference/kubectl/quick-reference/>
- Go over the CNCF Landscape
 - <https://landscape.cncf.io/>
- Practice Kubernetes Tutorials if you like
 - <https://kubernetes.io/docs/tutorials/>
- Review the KCNA Exam Resources
 - <https://training.linuxfoundation.org/certification/kubernetes-cloud-native-associate/>